# ECE391 Computer System Engineering Lecture 5

University of Illinois at Urbana- Champaign

Fall 2018

#### Lecture Topics

- Role of system software
- System calls, exceptions, & interrupts
- Processor support for interrupts

#### MP1 Handin and Demo Schedule

 Code must be committed to master branch in GitLab by 6PM on 9/17

- Handin Demo:
  - Monday 9/17, 6-8:30pm: Last name starts with A-K
  - Tuesday 9/18, 6-8:30pm: Last name starts with L-Z

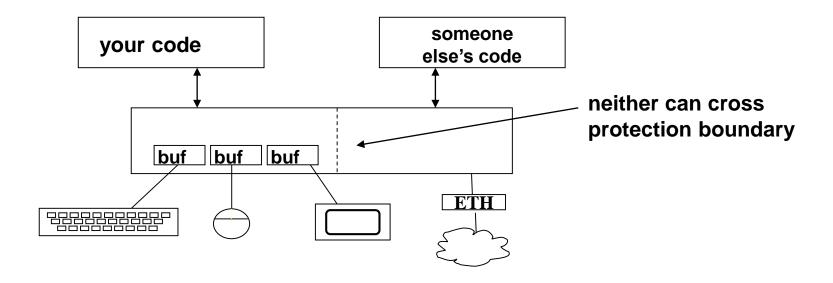
#### Role of System Software (1)

- System software serves three purposes
  - virtualization
  - protection
  - abstraction (particularly hiding asynchrony)
- virtualization:
  - the illusion of multiple/practically unlimited resources
- protection:
  - reduce/eliminate the chance of accidental and/or malicious destruction of data/results by another program

#### Role of System Software (2)

#### abstraction:

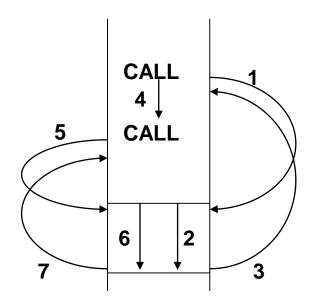
- hide fundamentally asynchronous nature of processor/device interaction
- provide simpler and more powerful interfaces (integrated w/protection)



#### System Calls, Interrupts, and Exceptions (1)

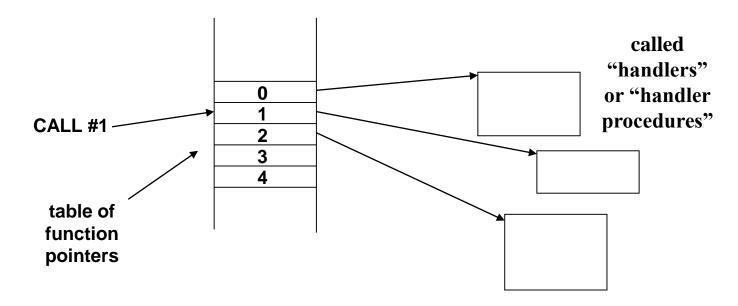
 recall that subroutines allow a programmer to encapsulate common operations

- the operating system
  - want to provide an interface including common operations
  - BUT don't want to re-link programs
  - NOR rely on everyone having exactly the same OS version



## System Calls, Interrupts, and Exceptions (2)

- solution
  - add a level of indirection!
- with indirection
  - to rewrite OS, just change the table
  - application code does not change



# System Calls, Interrupts, and Exceptions (3)

• in LC-3, we used the TRAP instruction; in x86, it's the INT instruction:

**INT** 8-bit imm. # (PUSH EIP), EIP ← table[imm8]

- the RTL is actually a little more complicated, as you'll see later in course
- called a trap (after instruction, or trap door through protection boundary)
- also called a system call (for operating system)

## System Calls, Interrupts, and Exceptions (4)

- vector tables/jump tables
  - i.e., tables of function pointers
  - convenient abstraction for many procedure-like activities
- Question:
  - What happens if software does something wrong, e.g.,
    - accesses a non-existent memory location?
    - issues an illegal/undefined instruction? divides by 0?
- What do we do to handle problems?
  - state machine that you design for processor may have don't cares
  - state machine that you build will do something (may be unknown)
  - so just let it run! (e.g., 6502 did so... and programmers used!)

#### System Calls, Interrupts, and Exceptions (5)

- a better solution: exceptions!
  - processor maps each problem to a vector #
  - calls procedure in vector table by #
- Where else might we use vector tables?

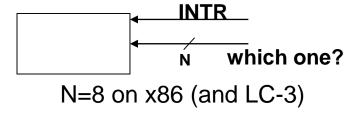
- Consider processor interactions with devices
  - a disk access takes about 10 milliseconds
  - new machines in lab: 10 ms = 32 million cycles
- should processor sit around asking, "Are my data here yet?"

#### System Calls, Interrupts, and Exceptions (6)

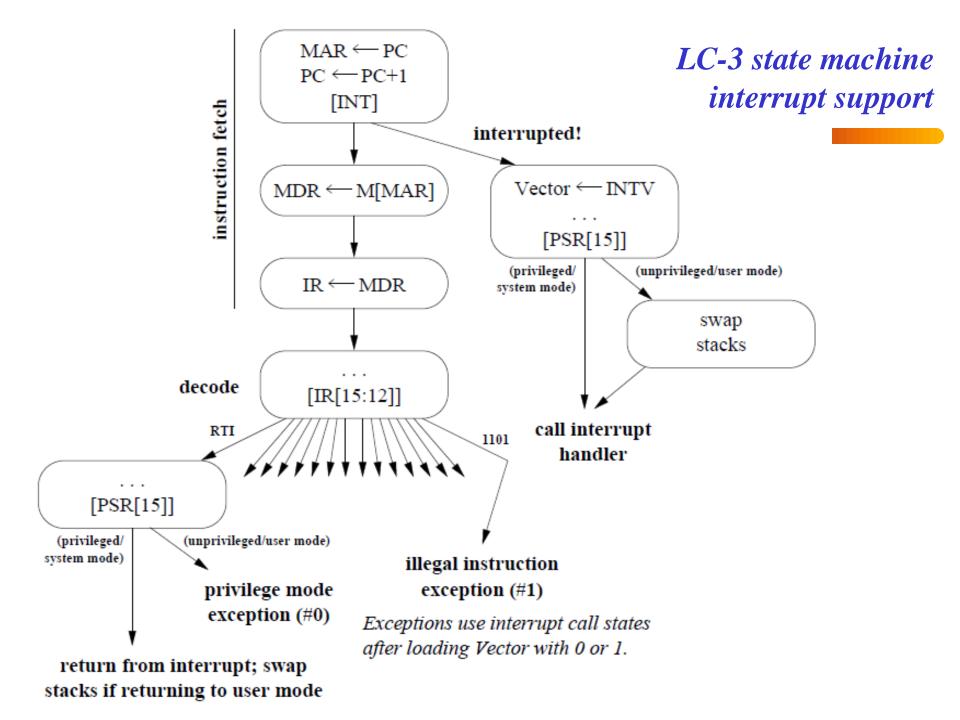
- analogous to posting a letter to a friend in Europe
- and checking your mailbox every minute for a reply
- instead, have your mail carrier ring your doorbell when it arrives
- in a processor, we call that an interrupt
- How can we use a vector table for interrupts?
  - each device has a vector #
  - call corresponding procedure in vector table when device generates
- x86 ISA
  - uses one table for all three kinds
  - called the Interrupt Descriptor Table (IDT)

#### Processor Support for Interrupts (1)

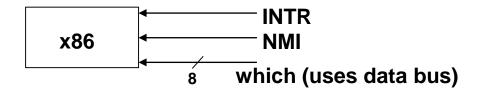
- How does a processor support interrupts?
- Logically...



 How should we change a processor's state machine to incorporate interrupts?



#### Processor Support for Interrupts (2)



- x86 allows software to block interrupts with a status flag (in EFLAGS)
- normal interrupts occur only when the interrupt enable flag (IF) is set
- some interrupts are too important
  - e.g., memory errors, power warnings, etc.
  - these are NOT maskable, and use a separate input to processor
  - called non-maskable interrupts (NMI)

## Interrupt Descriptor Table

- as mentioned earlier
  - x86 uses a single vector table
  - the Interrupt Descriptor Table (IDT)
  - hold vectors for interrupts, exceptions, and system calls
- note that this picture is partly OS-specific
  - the exception vector numbers are specified by Intel Why?
    - A: generated directly by processor's state machine
  - programmable interrupt controller (PIC) will be discussed later;
    - range of vectors generated is programmable, and is shown for Linux 2.4
  - note that a single entry is used for all system calls in Linux

	0x00	division error	
	0x02	NMI (	
	0x02	• • • • • • • • • • • • • • • • • • • •	
0x00-0x1F	0x03	breakpoint (used by KGDB) overflow	
0x00=0x1F	0x04	overnow	
	:		
defined	0x0B	segment not present	
by Intel	0x0C	stack segment fault	
	0x0D	general protection fault	
	0x0E	page fault	
	:		
	0x20	IRQ0 — timer chip	
	0x21		
0x20-0x27	0x22	IRQ2 — (cascade to slave)	
	ı	IRQ3	
master	0x24	IRQ4 — serial port (KGDB)	
8259 PIC	0x25	IRQ5	
	0x26	IRQ6	example
	1	IRQ7	of
	1	IRQ8 — real time clock	possible
	1	IRQ9	settings
0x28-0x2F	I	IRQ10	
		IRQ11 — eth0 (network)	
slave	1	IRQ12 — PS/2 mouse	
8259 PIC		IRQ13	
	1	IRQ14 — ide0 (hard drive)	
	0x2F	IRQ15	
0x30-0x7F	:	APIC vectors available to device drivers	
0x80	0x80	system call vector (INT 0x80)	
0x81-0xEE	:	more APIC vectors available to device drivers	
0xEF	0xEF	local APIC timer	
0xF0-0xFF	:	symmetric multiprocessor (SMP) communication vectors	

# Interrupt Descriptor Table

#### Device I/O

- How does a processor communicate with devices?
- Two possibilities
  - independent I/O use special instructions and a separate I/O port address space
  - memory-mapped I/O use loads/stores
     and dedicate part of the memory address space to I/O
- x86 originally used only independent I/O
  - but when used in PC, needed a good interface to video memory
  - solution? put card on the bus, claim memory addresses!
  - now uses both, although ports are somewhat deprecated

#### Device I/O

- I/O instructions have not evolved since 8086
  - 16-bit port space
    - byte addressable
    - little-endian (looks like memory)
  - instructions
    - IN port, dest.reg
    - OUT src.reg, port
  - the register operands are NOT general-purpose registers
    - all data to/from AL, AX, or EAX
    - port is either an 8-bit immediate (not 16!) or DX