# ECE391 Computer System Engineering Lecture 6

Dr. Zbigniew Kalbarczyk
University of Illinois at Urbana- Champaign

Fall 2018

# Lecture Topics

- System calls, exceptions, & interrupts (summary)
- Shared resources
- Critical sections

#### MP1 Handin and Demo Schedule

 Code must be committed to master branch in GitLab by 6PM on 9/17

- Handin Demo:
  - Monday 9/17, 6-8:30pm: Last name starts with A-K
  - Tuesday 9/18, 6-8:30pm: Last name starts with L-Z

# System Calls, Interrupts, and Exceptions (6)

#### x86 ISA

- uses one table for invocation of Interrupts, Exceptions, System Calls
- called the Interrupt Descriptor Table (IDT)

type	a synch ronous	unexpected	
interrupt	Yes	Yes	
coception	No	Tes	
system call	No	No	

	0x00	division error		
	002	NMI (		
	0x02 0x03	NMI (non-maskable interrupt)		
0x00-0x1F	0x03	,		
0x00=0x1F	0x04	overnow		
	:			
defined	0x0B	segment not present		
by Intel	0x0C	stack segment fault		
	0x0D	general protection fault		
	0x0E	page fault		
	:			
	0x20	IRQ0 — timer chip		
	0x21	IRQ1 — keyboard		
0x20-0x27	0x22	IRQ2 — (cascade to slave)		
	0x23	IRQ3		
master	0x24	IRQ4 — serial port (KGDB)		
8259 PIC	0x25	IRQ5		
	0x26		example	
		IRQ7	of	
	0x28		possible	
	0x29		settings	
0x28-0x2F	0x2A			
	0x2B			
slave	0x2C			
8259 PIC	0x2D	•		
	0x2E			
	0x2F	IRQ15		
0x30-0x7F	:	APIC vectors available to device drivers		
0x80	0x80	system call vector (INT 0x80)		
0x81-0xEE	:	more APIC vectors available to device drivers		
0xEF	0xEF	local APIC timer		
0xF0-0xFF	÷	symmetric multiprocessor (SMP) communication vectors		

# Interrupt Descriptor Table

# Shared Data and Resources (1)

- The question
  - interrupt handlers and programs share resources
    - What resources are shared between them?
    - How might interactions cause problems?
    - What can we do to fix those problems?

# Thought Problem on Shared Resources (2)

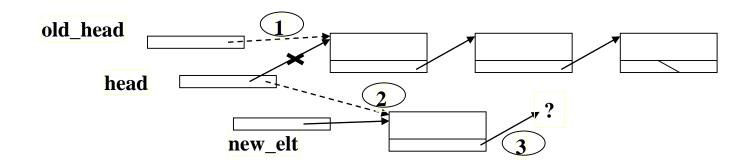
- Obvious things
  - registers
    - solution? save them to the stack
  - memory
    - solution? privatize
    - will still need to share some things; discussed later

# Thought Problem on Shared Resources (3)

#### Less obvious

- condition codes
  - solution? again, save them to the stack
- shared data
- More subtle
  - external state (e.g., on devices)
  - compiler optimization (e.g., volatility)
  - security leaks
    - e.g., application waits for interrupt, then observes values written by OS to stack
    - solution? use separate stack for kernel

## Example #1: a shared linked list



```
step 1: old_head = head;

step 2: head = new elt;

head = new_elt;
```

Oops! an interrupt!

# Examples of Shared Resources: Example #1: a shared linked list

#### The problem?

- linked list structure has invariant
- head points to first, chained through last via next field, ends with NULL
- complete operation maintains invariant
- partial operation does not need atomic update

# Examples of Shared Resources: Example #2: external state

#### The core problem

- devices have state
- processors interact with devices using specific protocols
- protocol often requires several steps (e.g., I/O instructions)
- device cannot differentiate which piece of code performed an operation

#### Example:

- VGA controller operations for scrolling window with color modulation
- interrupt handler drives color manipulations
- program handles scrolling using pixel shift
- both use VGA attribute register (port 0x3C0)

# Examples of Shared Resources: Example #2: external state

- Protocol for attribute control register
  - 22 different attributes accessed via this register
  - first send index
  - then send data
  - VGA tracks whether next byte sent is index or data

- Problem: processor can't know which one is expected
- Solution: reading from port 0x3DA forces VGA to expect index next

## Example #2: external state

- Consider the program code
  - the horizontal pixel panning register is register 0x13
  - assume that the code should write the value 0x03 to it

(discard)  $\leftarrow$  P[0x3DA] MOVW \$0x3DA, %DX INB (%DX),%AL

 $0x13 \rightarrow P[0x3C0]$  MOVW \$0x3C0, %DX

MOVB \$0x13, %AL

OUTB %AL, (%DX)

 $0x03 \rightarrow P[0x3C0]$  MOVB \$0x03, %AL

OUTB %AL, (%DX)

© Steven Lumetta, Zbigniew Kalbarczyk

# Examples of Shared Resources: Example #2: external state

- What happens if the interrupt occurs after the first write to 0x3C0?
  - the interrupt handler is executing basically the same code
  - leaves the VGA expecting an index

What is the solution?

## Example #3: handshake synchronization

- A device generates an interrupt after it finishes executing a command
- Consider the following attempt to synchronize

```
the shared variable...
```

```
int device is busy = 0;
```

the interrupt handler...

```
device is busy = 0;
```

## Example #3: handshake synchronization

The program function used to send a command to the device...

```
while (device_is_busy);/* wait until device is free */
device_is_busy = 1;
/* send new command to device */
```

- Q: Does the loop work?
- No.
  - Compiler assumes sequential program.
  - Variables can't change without code that changes them.

## Example #3: handshake synchronization

```
LOOP: MOVL device_is_busy, %EAX

CMPL $0, %EAX

JNE LOOP
```

- Nothing can change variable, so no need to reload (move LOOP down a line).
- Now nothing can change EAX, so move it down another line (to branch!).
- Will interrupt handler break you out of the resulting infinite loop?

# Examples of Shared Resources: Example #3: handshake synchronization

#### Solution

- mark variable as volatile volatile int device\_is\_busy;
- tells compiler to never assume that it hasn't changed between uses
- Why not mark everything volatile?
  - forces compiler to always re-load variables
  - more memory operations = slower program

# Examples of Shared Resources: Example #3: handshake synchronization

- Is it ok to swap setting the variable and sending the command?
- No.
  - introduces a race condition:

```
/* send new command to device */
---- INTERRUPT OCCURS HERE ----
device_is_busy = 1;
```

• Next command call blocks (forever) for device to be free.

# Examples of Shared Resources: Example #3: handshake synchronization

- Unfortunately, writing your code correctly is not enough.
  - compiler optimization allowed to reorder
    - so long as code is equivalent
    - assuming sequential program
  - also (and much more subtly!)
    - ISA implementation is allowed to reorder

#### Message

important to think about reordering possibilities by compiler and ISA and to prevent bad reorderings

#### Critical Sections

- Some parts of program need to appear to execute atomically, i.e., without interruption
- Full version: atomic with respect to code in interrupt handler
  - for now, the clause is implied i.e., only interrupt handlers can operate during our programs
  - however, multiprocessors may have >1 program executing at same time

#### Critical Sections

- Solution?
  - IF (the interrupt enable flag)
     critical section start (CLI)
     (the code to be executed atomically)
     critical section end (STI)
- What else must be prevented?
  - no moving memory ops into or out of critical section!

# Critical Sections in Examples

Example #2: external state

```
MOVW $0x3DA, %DX
                              CLI
       (%DX),%AL
INB
MOVW $0x3C0, %DX
                           the critical section
                           should be as short
MOVB $0x13, %AL
                              as possible
      %AL, (%DX)
OUTB
MOVB $0x03, %AL
      %AL, (%DX)
OUTB
                             STI
```

# Critical Sections in Examples

- Why should critical sections be short?
  - avoid delaying device service by interrupt handler
  - long delays can even crash system (e.g., swap disk driver timeout)
- Example #1: a shared linked list

```
old_head = head;
head = new_elt;
new_elt->next = old_head;
CLI

could skip first statement,
but including is safer

STI
```

- If interrupt handler can change list, too, leaving out first inst. creates race
- Example #3: handshake synchronization—volatile suffices for this example