# ECE391 Computer System Engineering Lecture 7

Dr. Zbigniew Kalbarczyk
University of Illinois at Urbana- Champaign

Fall 2018

#### Lecture Topics

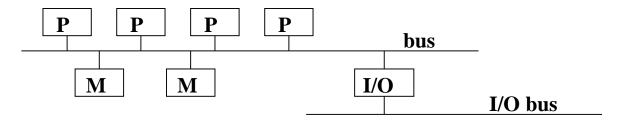
- Multiprocessors and locks
- Spin locks
- Synchronization issues

#### **Aministrivia**

PS2 Posted (soon)

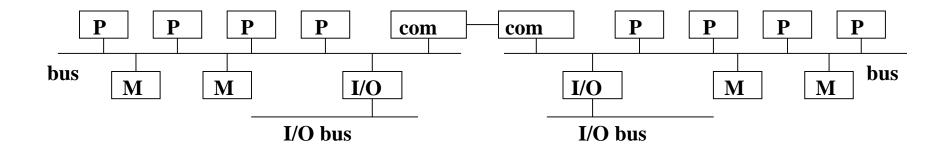
# Multiprocessors and Locks

- We solved the critical section problem for uniprocessors
- What about multiprocessors?
  - CLI ... critical section .... STI
- What is a multiprocessor?
  - usually a symmetric multiprocessor (SMP)



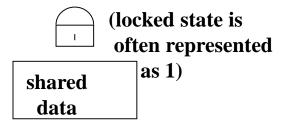
- symmetric aspect: all processors have equal latency to all memory banks
- multicore processors are similar from our perspective

Some non-uniform memory architecture (NUMA) machines were built



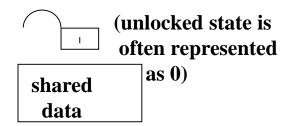
- Multithreaded code is not protected by IF
- Why haven't we solved the atomicity problem on multiprocessors?
  - interrupts are masked if *IF* is cleared!
  - answer: IF is not cleared on other processors!
  - just tell other processors to clear IF, too?
    - too slow
    - requires an interrupt!
- We need to use shared memory to synchronize...

- Logically, we use a lock
  - when we want to access a piece of shared data we first look at the lock
  - if it's locked, we wait until it's unlocked

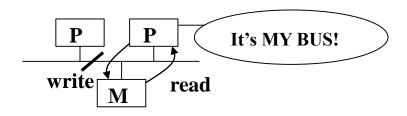




- we lock it
- access the data
- · then unlock it



 Locking must be atomic with respect to other processors!



read EAX
add 1

put it back
© Steven Lumetta, Zbigniew Kalbarczyk

LOCK: Prefix - execute following instruction with bus locked

LOCK INCL (%EAX)

#### Spin Locks

- The simplest lock
  - spin lock
  - lock op

```
do {
     try to change lock variable from 0 to 1
} while (attempt failed);
```

- other work to do? ignore it!
- spin in a tight loop on the lock (hence the name)
- Once successful, program/interrupt handler owns the lock

# Spin Locks (cont.)

- Only the owner can unlock
  - How?
    - (change lock variable to 0)
  - Need to be atomic?
    - (no, only owner can change when locked)
  - What about memory op reordering?
    - (must be avoided!)
- Why keep asking the last question?
  - reordering leads to subtle bugs
  - unlikely to happen often
  - unlikely to be repeatable
  - very hard to find!

# Linux Spin Lock API

Static initialization
 static spinlock\_t a\_lock = SPIN\_LOCK\_UNLOCKED;

Dynamic initialization

```
spin_lock_init (&a_lock);
```

- When is dynamic initialization safe?
  - lock must not be in use (race condition!)
  - other synchronization method must prevent use

#### Linux Spin Lock API – Basic Functions

```
void spin_lock (spinlock_t* lock);
void spin unlock (spinlock t* lock);
```

#### Linux Spin Lock API – Testing Functions

```
int spin_is_locked (spinlock_t* lock);
    returns 1 if held, 0 if not, but beware of races!
```

```
int spin_trylock (spinlock_t* lock);
```

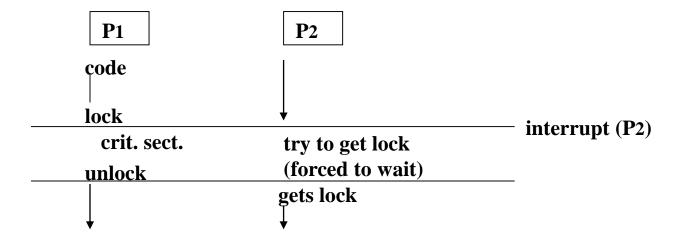
make one attempt; returns 1 on success, 0 on failure

```
void spin_unlock_wait (spinlock_t* lock);
```

wait until available (race condition again!)

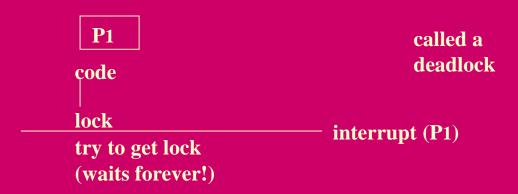
#### Linux Spin Lock API (cont.)

Is spinlock enough to protect a critical section?



#### Linux Spin Lock API (cont.)

What about ?



- Still need CLI/STI
- Which is first, CLI or lock?
  - CLI first
  - interrupt may occur between them, leading to scenario above

#### Linux' Lock/CLI Combo

```
static spinlock t the lock = SPIN LOCK UNLOCKED;
unsigned long flags;
spin lock irqsave (&the lock, flags); —
/* the critical section */
spin unlock irqrestore (&the lock, flags);
         asm volatile ("
              PUSHFL
             POPL %0
             CLI
         " : "=g" (flags) /* outputs */
                      /* inputs */
           : "memory" /* clobbers */
         spin lock (&the lock);
© Steven Lumetta, Zbigniew Kalbarczyk
```

#### Linux' Lock/CLI Combo (cont)

```
spin_unlock (&the_lock);
asm volatile ("
    PUSHL %0
    POPFL
                   /* outputs */
  : "g" (flags) /* inputs */
  : "memory", "cc" /* clobbers */
```

#### Comments on code

- Notice that spin\_lock\_irqsave changes the flags argument (it's a macro)
- The "memory" argument
  - tells compiler that all memory is written by assembly block
  - prevents compiler from moving memory ops across assembly
- The "cc" argument
  - condition codes change
  - can lead to subtle bugs if left out!
- spin\_lock and spin\_unlock calls
  - become NOPs on uniprocessors
  - in that case, calls just change IF
- Restore rationale: may have had IF=0 on entry; if so, STI is unsafe

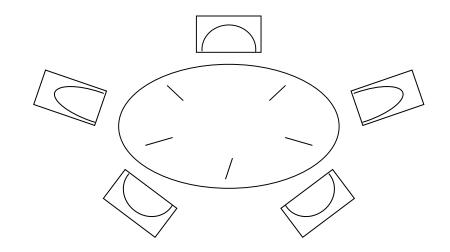
#### Another Philosophy Lesson

#### Synchronization issues

- five hungry philosophers
- five chopsticks

#### Protocol

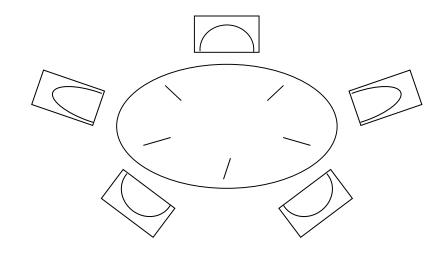
- take left chopstick (or wait)
- take right chopstick (or wait)
- eat
- release right chopstick
- release left chopstick
- digest
- repeat



#### Another Philosophy Lesson (cont.)

- How about the following protocol?
  - take left chopstick (or wait)
  - if right chopstick is free, take it
  - else release left chopstick and start over
  - eat
  - release right
  - release left
  - digest
  - repeat





#### Another Philosophy Lesson (cont.)

What if all philosophers act in lock-step (same speed)?

left left left left left release release release release left left left left left release release release release left left left left left (ad infinitum)

Called a livelock

#### Another Philosophy Lesson (cont.)

- To solve the problem, need (partial) lock ordering
  - e.g., call chopsticks #1 through #5
  - protocol: take lower-numbered, then take higher-numbered
  - two philosophers try to get #1 first
  - can't form a loop of waiting philosophers
  - thus someone will be able to eat

# Example of Code Understanding

Given code in left column

– Which other columns are compatible (i.e., can't lead to deadlock)?

```
#0 #1 #2 #3
lock A lock B
lock B lock A lock A lock B
. . . . .
. . . . .
. unlock B unlock A unlock A unlock B
unlock A
```

Columns 2 and 3 are compatible with column 0