# ECE391 Computer System Engineering Lecture 8

Dr. Zbigniew Kalbarczyk
University of Illinois at Urbana- Champaign

Fall 2018

# Lecture Topics

- Synch issues
- Conservative synchronization design
- Semaphores
- Reader/writer synchronization
- Selecting a synchronization mechanism

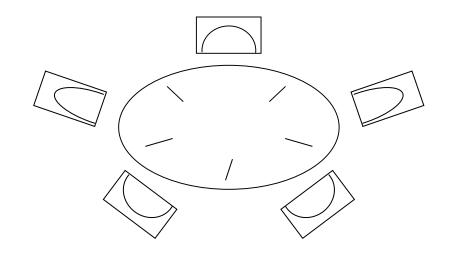
# Another Philosophy Lesson

#### Synchronization issues

- five hungry philosophers
- five chopsticks

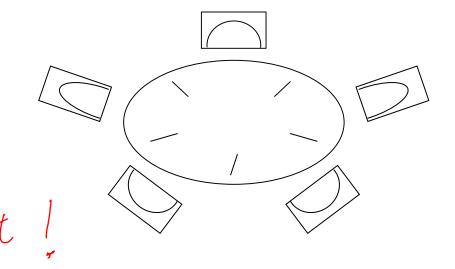
#### Protocol

- take left chopstick (or wait)
- take right chopstick (or wait)
- eat
- release right chopstick
- release left chopstick
- digest
- repeat



# Another Philosophy Lesson (cont.)

- How about the following protocol?
  - take left chopstick (or wait)
  - if right chopstick is free, take it
  - else release left chopstick and start over
  - eat
  - release right
  - release left
  - digest
  - repeat
- Does this work?



## Another Philosophy Lesson (cont.)

What if all philosophers act in lock-step (same speed)?

```
left
        left left
                       left
                              left
release release release release
 left
        left left
                       left
                              left
release release release release
 left
       left
                left
                       left
                              left
         (ad infinitum)
```

Called a livelock

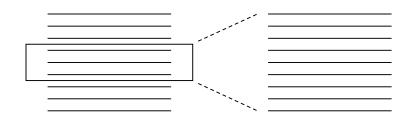
## Another Philosophy Lesson (cont.)

- To solve the problem, need (partial) lock ordering
  - e.g., call chopsticks #1 through #5
  - protocol: take lower-numbered, then take higher-numbered
  - two philosophers try to get #1 first
  - can't form a loop of waiting philosophers
  - thus someone will be able to eat

# Conservative Synchronization Design

- Getting synchronization correct can be hard
  - it's the focus of several research communities

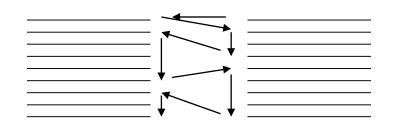
- On uniprocessor
  - mentally insert handler code between every pair of adjacent instructions
  - ask whether anything bad can happen
  - if so, prevent with CLI/STI



# Conservative Synchronization Design (cont)

#### On a multiprocessor

- consider all possible interleavings of instructions
- amongst all pieces of (possibly concurrent) code
- ask whether anything bad can happen
- if so, use a lock to force serialization
- good luck!



# Conservative Synchronization Design (cont)

What does "bad" mean, anyway?

A <u>conservative</u> but <u>systematic</u> definition

if any data written by one piece of code

are also read or written by another piece of code

these two pieces <u>must be atomic</u> with respect to each other

# Conservative Synchronization Design (cont)

- What variables are shared?
- step 0: ignore the parts that don't touch shared data
- step 1: calculate read & write sets
- step 2: check for R/W, W/W relationships

#### – must be atomic!

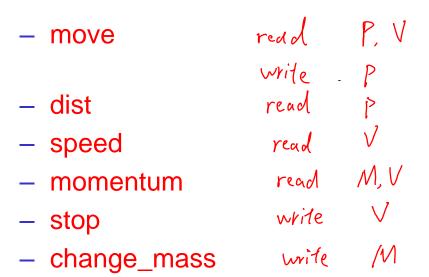
- step 3: add lock(s) to guarantee atomic execution (pick order if > 1 locks)
- step 4: optimize if desired

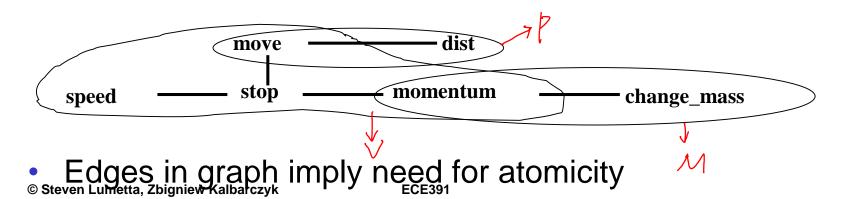
```
typedef struct {
    double mass:
    double x, y, z; /* position */
    double vx, vy, vz; /* velocity */
} thing t;
void move (thing t* t)
    t \rightarrow x += t \rightarrow vx;
   t->y += t->vy;
    t->z += t->vz;
double dist(thing t* t)
    return sqrt(t->x * t->x +
                t->y * t->y +
                t->z * t->z);
double speed(thing_t* t)
    return sqrt(t->vx * t->vx +
                t->vy * t->vy +
                t->vz * t->vz);
double momentum(thing t* t)
    double tmp = t->mass;
    return tmp * speed(t);
void stop(thing t* t)
    t->vx = t->vy = t->vz = 0;
void change mass(thing t* t, double new mass)
    t->mass = new mass;
```

```
#include<stdio.h>
typedef struct person t person t;
struct person t {
    char*
               name;
    int
               age;
    person t* next;
};
static person t* group;
void birthday(person t* p)
    p->age++;
void list people(void)
    person t* p;
    int num;
    for (p=group, num=0; NULL != p; p=p->next, num++)
        if (10 > num)
            printf("%s %d\n", p->name, p->age);
            printf("%s\n", p->name);
```

# Conservative Synchronization Design – Example Code Analysis

#### Read/write sets





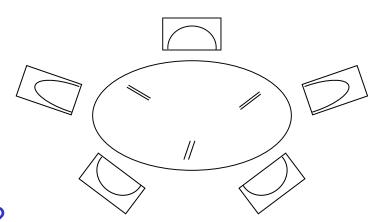
# Conservative Synchronization Design – Example Code Analysis (cont)

- Each lock = circle in graph
- All edges must be contained in some circle

- One lock suffices, but prevents parallelism (performance)
- Could use three (as shown above);
   then MUST pick a lock order!

# Role of Semaphores

- Recall our philosophical friends
- Five philosophers
- Three pairs of chopsticks (one "lock" per pair)
- Problem: how do you get a pair?

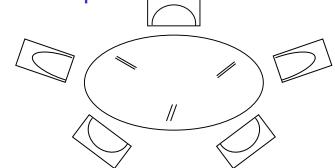


- Option 1: walk around the table until you find a pair free
  - lots of walking
  - other people may cut in front of you

# Role of Semaphores

- Option 2: pick a target pair and wait for it to be free
  - other pairs may be on the table
  - but you're still waiting hungrily for your chosen pair

- Instead, use a semaphore!
  - an atomic counter
  - proberen (P for short, meaning test/decrement)
  - verhogen (V for short, meaning increment)
  - Dutch courtesy of E. Dijkstra



# **Semaphores**

- When are semaphores useful?
- Fixed number of resources to be allocated dynamically
  - physical devices
  - virtual devices
  - entries in a static array
  - logical channels in a network
- Linux semaphores have a critical difference from Linux spin locks
  - can block (sleep) and let other programs run (spin locks do not block)
  - thus <u>must not</u> be used by interrupt handlers
  - used for system calls, particularly with long critical sections

## Linux Semaphore API

```
void sema init (struct semaphore* sem, int val);
Initialize dynamically to a specified value
void init MUTEX (struct semaphore* sem);
Initialize dynamically to value one (mutual exclusion)
void down (struct semaphore* sem);
Wait on the semaphore (P)
void up (struct semaphore* sem);
Signal the semaphore (V)
```

# Linux Semaphore API (cont.)

- If critical section needs both semaphores and spin locks
  - Get all semaphores first!
  - Linux expects <u>not</u> to be preempted while holding spin locks
  - Semaphore code <u>voluntarily</u> relinquishes processor

# Reader/Writer Problem: The Philosophers and Their Newspaper

- Philosophers like to read newspaper
  - each philosopher reads frequently, taking short breaks between
  - multiple philosophers may read at same time
    - different sections
    - or just over another's shoulder
- Paper carrier delivers new paper
  - once per day (infrequently)
  - must change all sections at once
- Reader/writer synchronization supports this style
  - allows many (in theory infinite) readers simultaneously
  - at most one writer (and never at same time as readers)

# Reader/Writer Problem: The Philosophers and Their Newspaper (cont)

- What if newspaper is always being read? starvation!
- Linux provides two types of reader/writer synchronization
  - reader/writer spin locks (rwlock\_t)
    - extension of spin locks
    - use for short critical sections
    - ok to use in interrupt handlers
    - admit writer starvation (you must ensure that this not happen)
  - reader/writer semaphores (struct rw\_semaphore)
    - extension of semaphores
    - use only in system calls
    - do not admit writer starvation (new readers wait if writer is waiting)

## Selecting a Synchronization Mechanism

- General questions for synchronization of shared resources
  - What are the operations?
  - What are their frequencies?
  - example:
    - walk list and find min. age?
    - update min. age every time list changes?
  - What are the shared data?
- Goals
  - short critical sections
  - low contention for data

# Selecting a Synchronization Mechanism (cont)

Safe selection (when in doubt...)

	mutual exclusion	reader/writer
data shared by interrupt handlers	spin_lock_irqsave spin_unlock_irqrestore	read/writelock_irqsave read/write_unlock_irqrestore
data shared only by system calls	down up	down_read/write up_read/write