ECE391 Computer System Engineering Lecture 9

Dr. Zbigniew Kalbarczyk
University of Illinois at Urbana- Champaign

Fall 2018

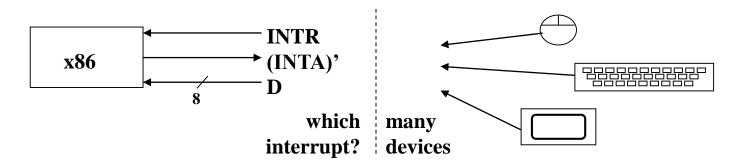
Lecture Topics

- Programmable interrupt controller (PIC)
 - motivation & design
 - hardware for x86
 - Linux abstraction of PIC

ECE391 EXAM 1

- EXAM I October 2; 7:00pm-9:00pm
 - Seating: TBA
- Conflict Exam
 - Tuesday October 2, Time: TBA
- NO Lecture on Tuesday, October 2
- Review Session
 - Saturday September 29, 4:00pm
 - Location: ECEB 1002

PIC Motivation and Design



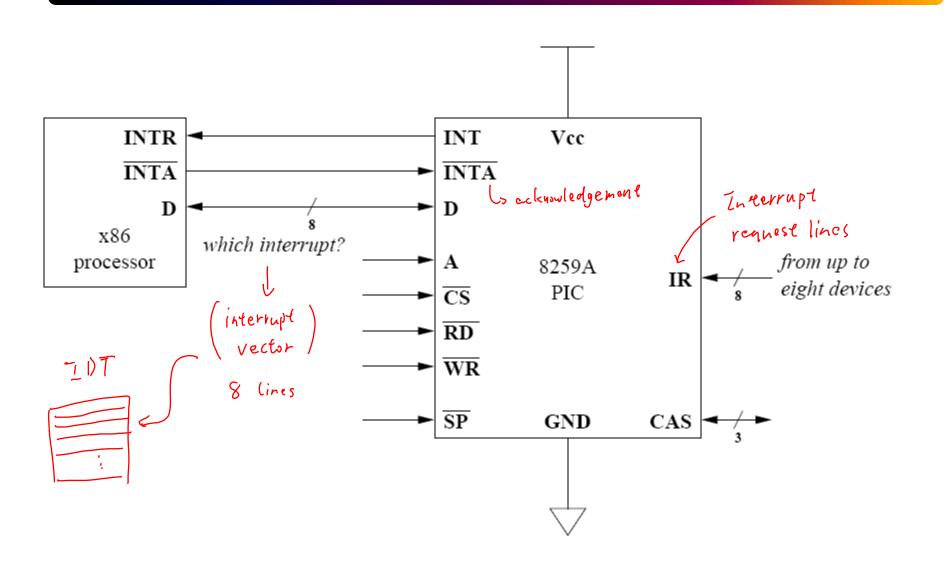
- How do we connect devices to the processor's interrupt input?
- An OR gate? why not?
 - who writes the vector #?
 - possible to build arbiter, but...
 - what if more than one raised interrupt?
 - extra work for processor to query all devices
 - must execute interrupt code for device that raised interrupt
 - could have been more than one device
 - no way to tell with OR gate

PIC Motivation and Design (cont.)

- not all devices support query
 - many devices too simplistic to support query
 - and operations (e.g., reading from port) may not be idempotent
- nice to have concept of priority and preemption,
 i.e., interrupting an interrupt handler

8259A

Programmable Interrupt Controller (PIC)



Logical Model of PIC Behavior

- Watch for interrupt signals
 - from up to eight devices
 - one interrupt line each

- Using internal state
 - track which devices/input lines
 - are currently in service by processor
 - i.e., processor is executing interrupt handler for that interrupt

- When a device raises an interrupt
 - check if priority is higher than those of in-service interrupts
 - if not, do nothing
 - if so
 - report the highest-priority raised to the processor
 - mark that device as being in service

- When processor reports EOI (end of interrupt) for some interrupt
 - remove the interrupt from the in-service mask
 - check for raised interrupt lines
 - that should be reported to processor

- Protocol for reporting interrupts
 - PIC raises INTR
 - processor strobes INTA' (active low) repeatedly
 - creates cycles for PIC to write vector to data bus
 - (must follow spec timing! PIC is not infinitely fast!)
 - processor sends EOI with specific combinations of A & D inputs (A is from address bus, D is from data bus)

- What about A, CS', RD', and WR'?
 - A = address match for ports
 - CS' = chip select (does processor want PIC to read/write?)
 - RD' and WR' defined from processor's point of view
 - RD' = processor will read data (vector #) from PIC
 - WR' = processor will write data (command, EOI) to PIC

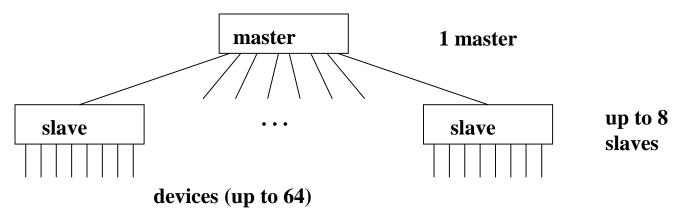
is Acrive Low

- Map to ports 0x20 & 0x21
 - given ADDR bus
 - logic to form A and CS' inputs to PIC

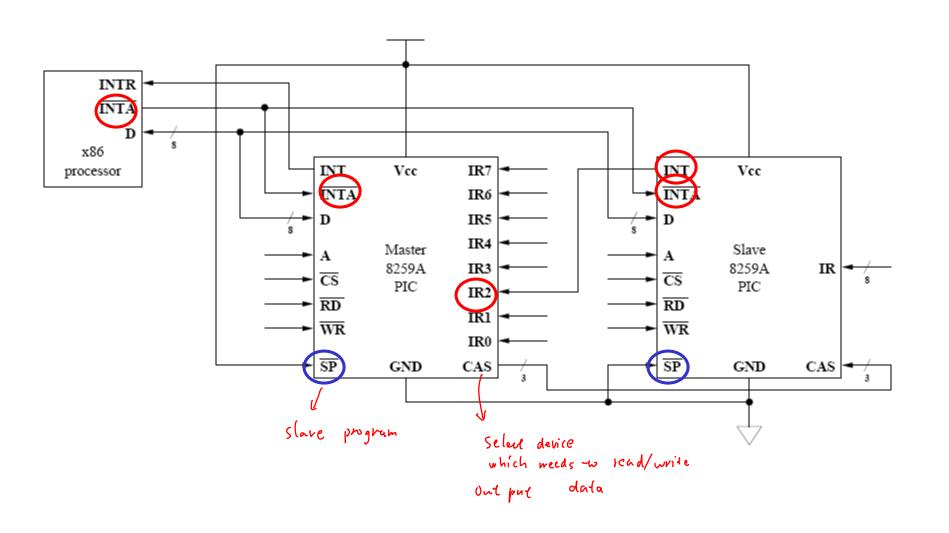
ADDR bus ADDR[0] ADDR[15:1] = 0x10to CS' to A on PIC

- Remember
 - the PIC is asynchronous!
 - all interactions have timing constraints
 - see specifications for details

- What about SP' & CAS?
- Are eight devices enough? No? What then?
 - hook 2, 3, or 4 PICs to processor?[same problem as before!]
 - design a big PIC?[waste of transistors; not cheap in 8259A era]
 - design several PICs?[waste of humans! (and production costs)]
- Better answer: cascade



Cascade Configuration of PICs



PIC (cont.)

- Previous figure showed x86 configuration of two 8259A's
 - master 8259A mapped to ports 0x20 & 0x21
 - slave 8259A mapped to ports 0xA0 & 0xA1
 - slave connects to IR2 on master

Question

- prioritization on 8259A: 0 is high, 7 is low
- what is prioritization across all 15 pins in x86 layout?
- (highest) M0...M1...S0...S7...M3...M7 (lowest)

PIC (cont.)

- In Linux (initialization code to be seen shortly)
 - master IR's mapped to vector #'s 0x20 0x27
 - slave IR's mapped to vector #'s 0x28 0x2F
 - remember the IDT?

	0x00	division error	
	0x02	NMI (
	0x02	• • • • • • • • • • • • • • • • • • • •	
0x00-0x1F	0x03	breakpoint (used by KGDB) overflow	
0x00=0x1F	0x04	overnow	
	:		
defined	0x0B	segment not present	
by Intel	0x0C	stack segment fault	
	0x0D	general protection fault	
	0x0E	page fault	
	:		
	0x20	IRQ0 — timer chip	
	0x21		
0x20-0x27	0x22	IRQ2 — (cascade to slave)	
	ı	IRQ3	
master	0x24	IRQ4 — serial port (KGDB)	
8259 PIC	0x25	IRQ5	
	0x26	IRQ6	example
	1	IRQ7	of
	1	IRQ8 — real time clock	possible
	1	IRQ9	settings
0x28-0x2F	I	IRQ10	
		IRQ11 — eth0 (network)	
slave	1	IRQ12 — PS/2 mouse	
8259 PIC		IRQ13	
	1	IRQ14 — ide0 (hard drive)	
	0x2F	IRQ15	
0x30-0x7F	:	APIC vectors available to device drivers	
0x80	0x80	system call vector (INT 0x80)	
0x81-0xEE	:	more APIC vectors available to device drivers	
0xEF	0xEF	local APIC timer	
0xF0-0xFF	:	symmetric multiprocessor (SMP) communication vectors	

Interrupt Descriptor Table

Linux 8259A Initialization

```
void init 8259A(int auto eoi)
        unsigned long flags;
        i8259A auto eoi = auto eoi;
        spin lock irqsave(&i8259A lock, flags)
        outb(0xff, 0x21);
                                /* mask all of 8259A-1 */
        outb(0xff, 0xA1);
                                /* mask all of 8259A-2 */
         * outb p - this has to work on a wide range of PC hardware.
        outb p(0x11, 0x20);
                                /* ICW1: select 8259A-1 init */
        outb p(0x20 + 0, 0x21); /* ICW2: 8259A-1 IRO-7 mapped to 0x20-0x27 */
                              /* 8259A-1 (the master) has a slave on IR2 */
        outb p(0x04, 0x21);
        if (auto eoi)
                                        /* master does Auto EOI */
                outb p(0x03, 0x21);
        else
                outb p(0x01, 0x21);
                                        /* master expects normal EOI */
        outb p(0x11, 0xA0);
                              /* ICW1: select 8259A-2 init */
        outb p(0x20 + 8, 0xA1); /* ICW2: 8259A-2 IRO-7 mapped to 0x28-0x2f */
        outb_p(0x02, 0xA1);
                              /* 8259A-2 is a slave on master's IR2 */
        outb p(0x01, 0xA1);
                                /* (slave's support for AEOI in flat mode
                                    is to be investigated) */
        if (auto eoi)
                 * in AEOI mode we just have to mask the interrupt
                 * when acking.
                i8259A irq type.ack = disable 8259A irq;
        else
                i8259A irq type.ack = mask and ack 8259A;
        udelay(100);
                                /* wait for 8259A to initialize */
        outb(cached 21, 0x21); /* restore master IRQ mask */
        outb(cached A1, 0xA1); /* restore slave IRQ mask */
        spin unlock irqrestore(&i8259A lock, flags);
```

Comments on Linux' 8259A Initialization Code

- 1. What is the auto_eoi parameter? always = 0
- 2. Four initialization control words to set up the master 8259A

info contained in Initialization Control Word

3. Four initialization control words to set up the slave 8259A

```
ICW1 0 start init, edge-triggered inputs, cascade mode, 4 ICWs ICW2 1 high bits of vector #
```

ICW3 1 master: bit vector of slaves; slave: input pin on master

ICW4 1 ISA=x86, normal/auto EOI

port(A=?)

Comments on Linux' 8259A Initialization Code (cont.)

- 5. What does the "_p" mean on the "outb" macros?
 - add PAUSE instruction after OUTB; "REP NOP" prior to P4
 - delay needed for old devices that cannot handle processor's output rate
- 6. Critical section spans the whole function; why?
 - avoid other 8259A interactions during initialization sequence
 - (device protocol requires that four words be sent in order)
- 7. Why use _irqsave for critical section?
 - this code called from other interrupt initialization routines
 - which may or may not have cleared IF on processor