ECE391 Computer System Engineering Lecture 13

Dr. Zbigniew Kalbarczyk
University of Illinois at Urbana- Champaign

Fall 2018

Lecture Topics

- Soft interrupts (tasklets) in Linux
- Summary of Linux interrupt system

Aministrivia

- MP3 Teams
 - Due by end of Wednesday, October 10

Soft Interrupts in Linux linux/interrupt.h and kernel/softirq.c

- When are soft interrupts executed?
 - after a hard interrupt completes
 - periodically by a daemon in the kernel

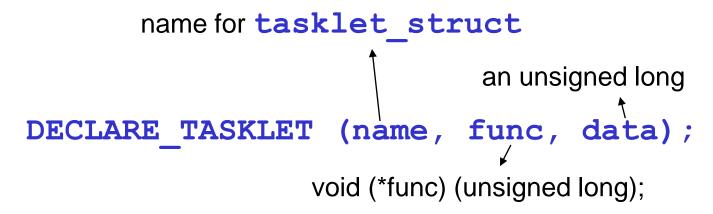
How?

- seven or eight prioritized types, including high and low tasklet priorities
- linked list for each tasklet priority
- run on processor on which interrupt is scheduled
- each handler atomic with respect to itself (only)

Soft Interrupts in Linux (cont.)

linux/interrupt.h and kernel/softirq.c

Declaring a handler



next -	J linked list
state	TASKLET_STATE_SCHED, TASKLET_STATE_RUN
count	# of disables
func	pointer to the tasklet function
data	integer which can be used by the tasklet function

Tasklet Scheduling

The following two calls schedule a tasklet for execution

```
void tasklet_schedule (struct tasklet_struct* t);
void tasklet_hi_schedule (struct tasklet_struct* t);
```

- First form
 - schedules tasklet at low priority
 - on the executing processor
- Second form schedules at high priority
- Enable and disable calls analogous to hard interrupts (including nesting)

Tasklet Execution

- do_softirq call
 - checks per-processor bit vector of pending priorities (high, low, etc.)
 - executes action for each priority [softirq_vec]
 - tasklet_action walks through linked list
 [tasklet_hi_action walks through high-priority list]
 - repeats up to 10 times or until no softirqs are raised

Tasklet Execution Atomicity

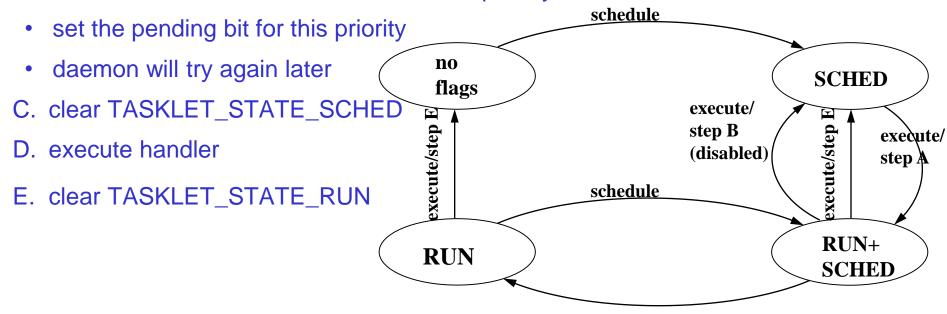
- Two bits in state changed atomically
 - TASKLET_STATE_SCHED tasklet scheduled for execution
 - TASKLET_STATE_RUN tasklet executing on some processor

- When scheduling
 - set TASKLET_STATE_SCHED atomically
 - if already set, schedule call does nothing

Tasklet Execution Atomicity (cont.)

When executing, for each tasklet in linked list (at either priority)

- A. set TASKLET_STATE_RUN atomically (if already set, stop)
- B. check if tasklet is software disabled (count field)
 - if so, clear TASKLET_STATE_RUN
 - leave the tasklet in the linked list for this priority



execute/step C execution (execute/step D) occurs in lower two states (and execute/step A fails in these states)

Interrupt Descriptor Table (IDT)

- Associates the interrupt line with the int. handler routine.
- 256 entries (each 8-bytes) or descriptors; each corresponds to an interrupt vector
 - hardware interrupts mapped into vectors 0x20 to 0x2F
- All Linux interrupt handlers are activated by so called: interrupt gates (a descriptor type)
- Whenever an interrupt gate is hit, interrupts are disabled automatically by the processor

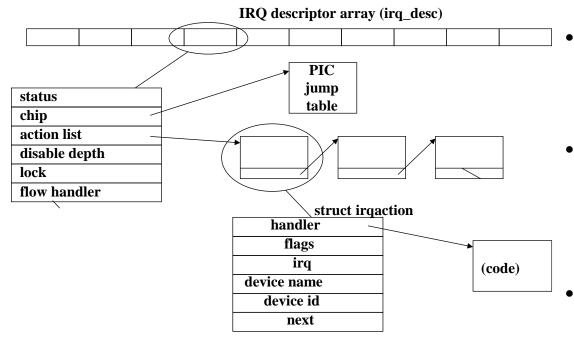
IDT
interrupt[x]

 Before the kernel enables interrupts it must initialize the *idtr* register to point to the IDT table (set by the kernel using lidt instruction)

IDT Initialization

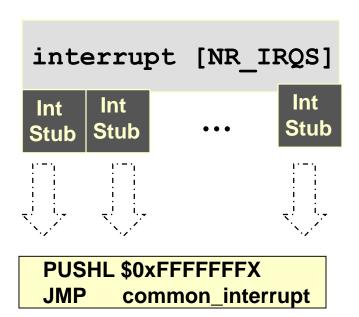
Kernel initialization
 (setup_idt()) fills
 all the 256 entries of
 IDT with the
 provisional (or null)
 handler

irq_descArray



- Every interrupt vector has its own irq_desc_t descriptor
- Descriptors are grouped together in irq_desc array, a data structure supported by Linux
- When a device driver calls
 the request_irq()
 function a new structure to
 represent the handler is
 allocated and initialized

interrupt[NR_IRQS] Array



 Kernel maintains one global array of function pointers
 (interrupt [NR_IRQS]) in which it stores pointers to interrupt stubs
 (NR_IRQS is 16 if we use the PIC)

Initialization of Interrupt Gates

- During initialization init_IRQ() sets the status field of each IRQ descriptor to IRQ_DISABLED
- init_IRQ() updates the IDT by replacing the provisional interrupt gates with new ones

```
for (i = 0; i < NR_IRQS; i++)
    if (i+32 != 128)
        set_intr_gate(i+32, interrupt[i]);</pre>
```

 Interrupt gates are set to the addresses found in the interrupt [NR_IRQS] array