ECE391
Computer System Engineering

Lecture 19

Dr. Jian Huang

University of Illinois at Urbana- Champaign

Fall 2018

ECE391 EXAM 1

- **EXAM II – November 6; 7:00pm-9:00pm**
  - Seating:
    - ECEB 1002: last name starting with A to W
    - ECEB 1013: last name starting with X to Z
- **Review Session**
  - In Class; Thursday November 1
- **Conflict Exam**
  - Tuesday November 6, 2:00pm; Location: CSL 141
  - By Friday November 2 send email to kalbarcz@Illinois.edu to request Conflict Exam
- **NO Lecture on Tuesday, November 6**

Lecture Topics

- System Calls

System Calls (Definition & An Example)

- Unix/Linux uses systems calls to implement most interfaces between User Mode processes and hardware device

➢ **write (fd, &buf, len)**

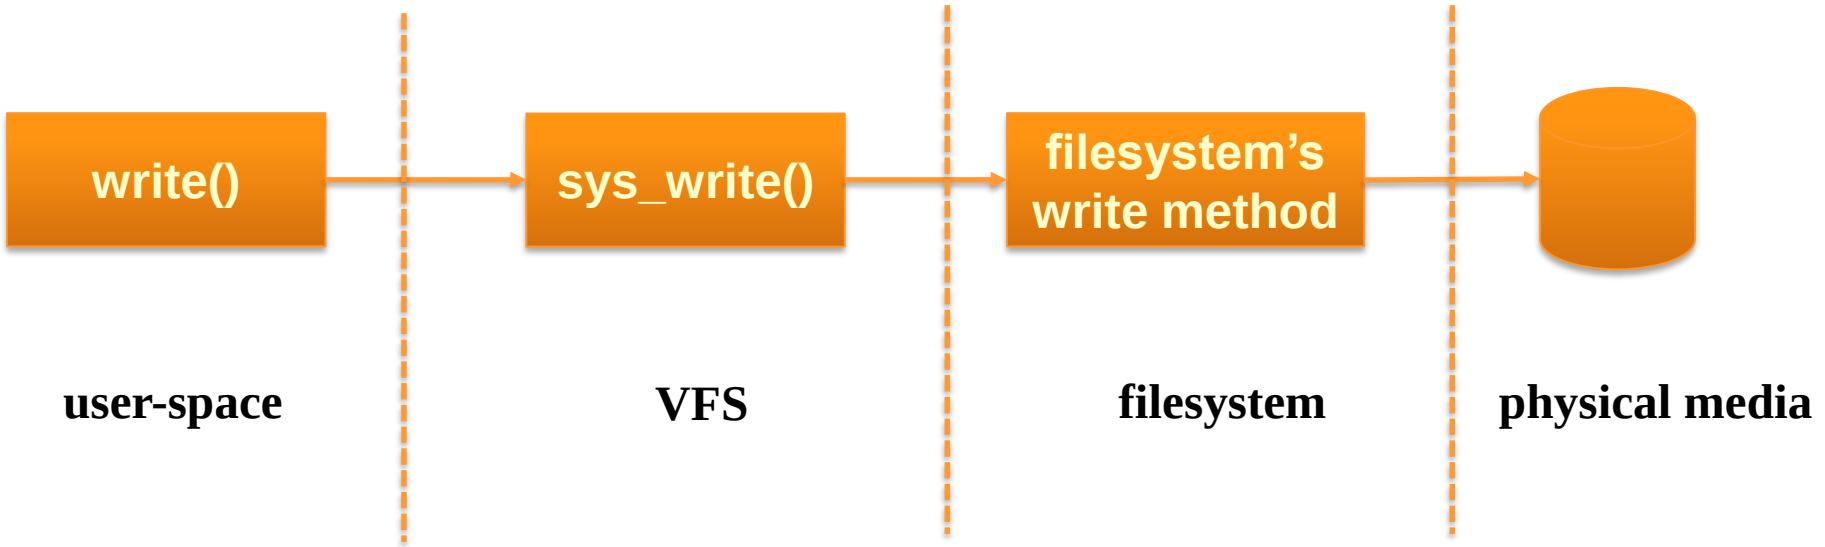

Why System Call ?

- Making programming easier (hiding low-level hardware)
- Increasing system security (checking the correctness)
- Making programs more portable (the same interfaces for different kernel)

System Call Categories

- Process management
- Memory management
- File management
- Device management
- Communication
- ……

User Mode vs. Kernel Mode

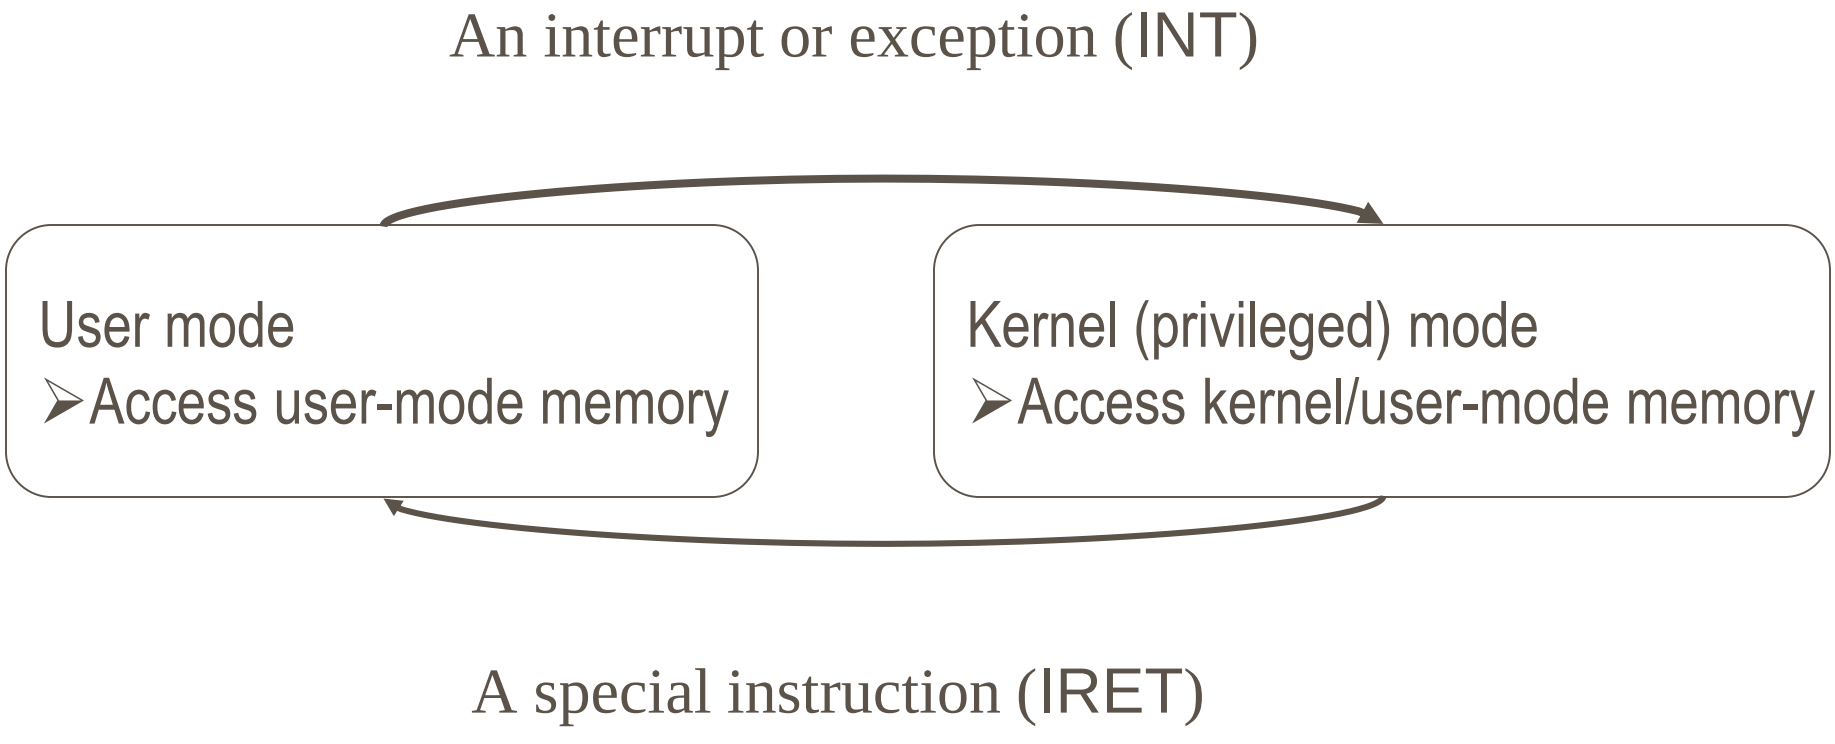

Why Privilege Mode ?

- Special Instructions
  - Mapping, TLB, etc
  - Device registers
  - I/O channels, etc.
- Processor Features
  - SSE3 (Streaming Single-Instruction-Multiple-Data Extension 3)
- Device access

X86 Protection Ring

Privileged instructions
Can be executed only
When current privileged
Level (CPR) is 0

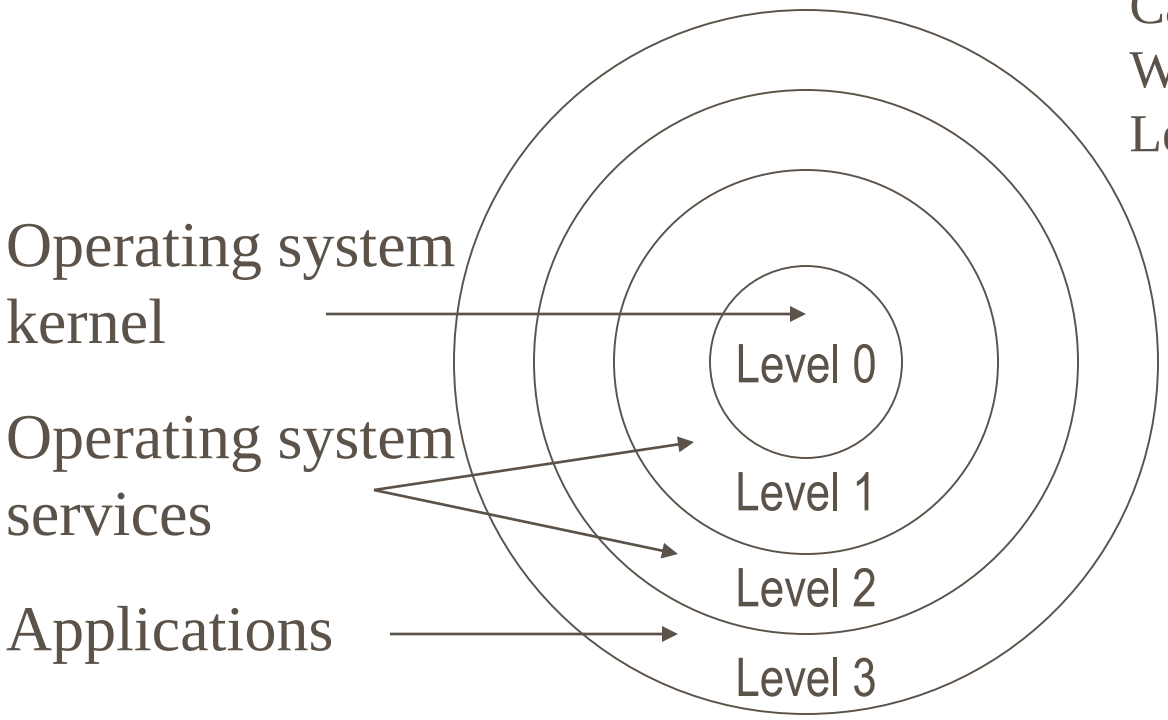

System Call Mechanism

- User code can be arbitrary
- User code cannot modify kernel memory
- Makes a system call with parameters
- The call mechanism switches code to kernel mode
- Execute system call
- Return with results

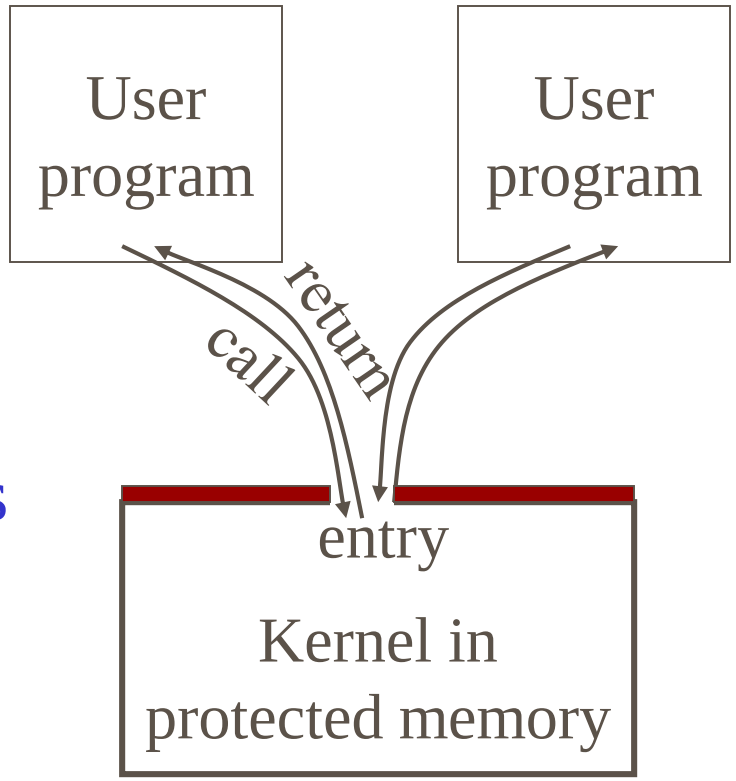

System Call Implementation

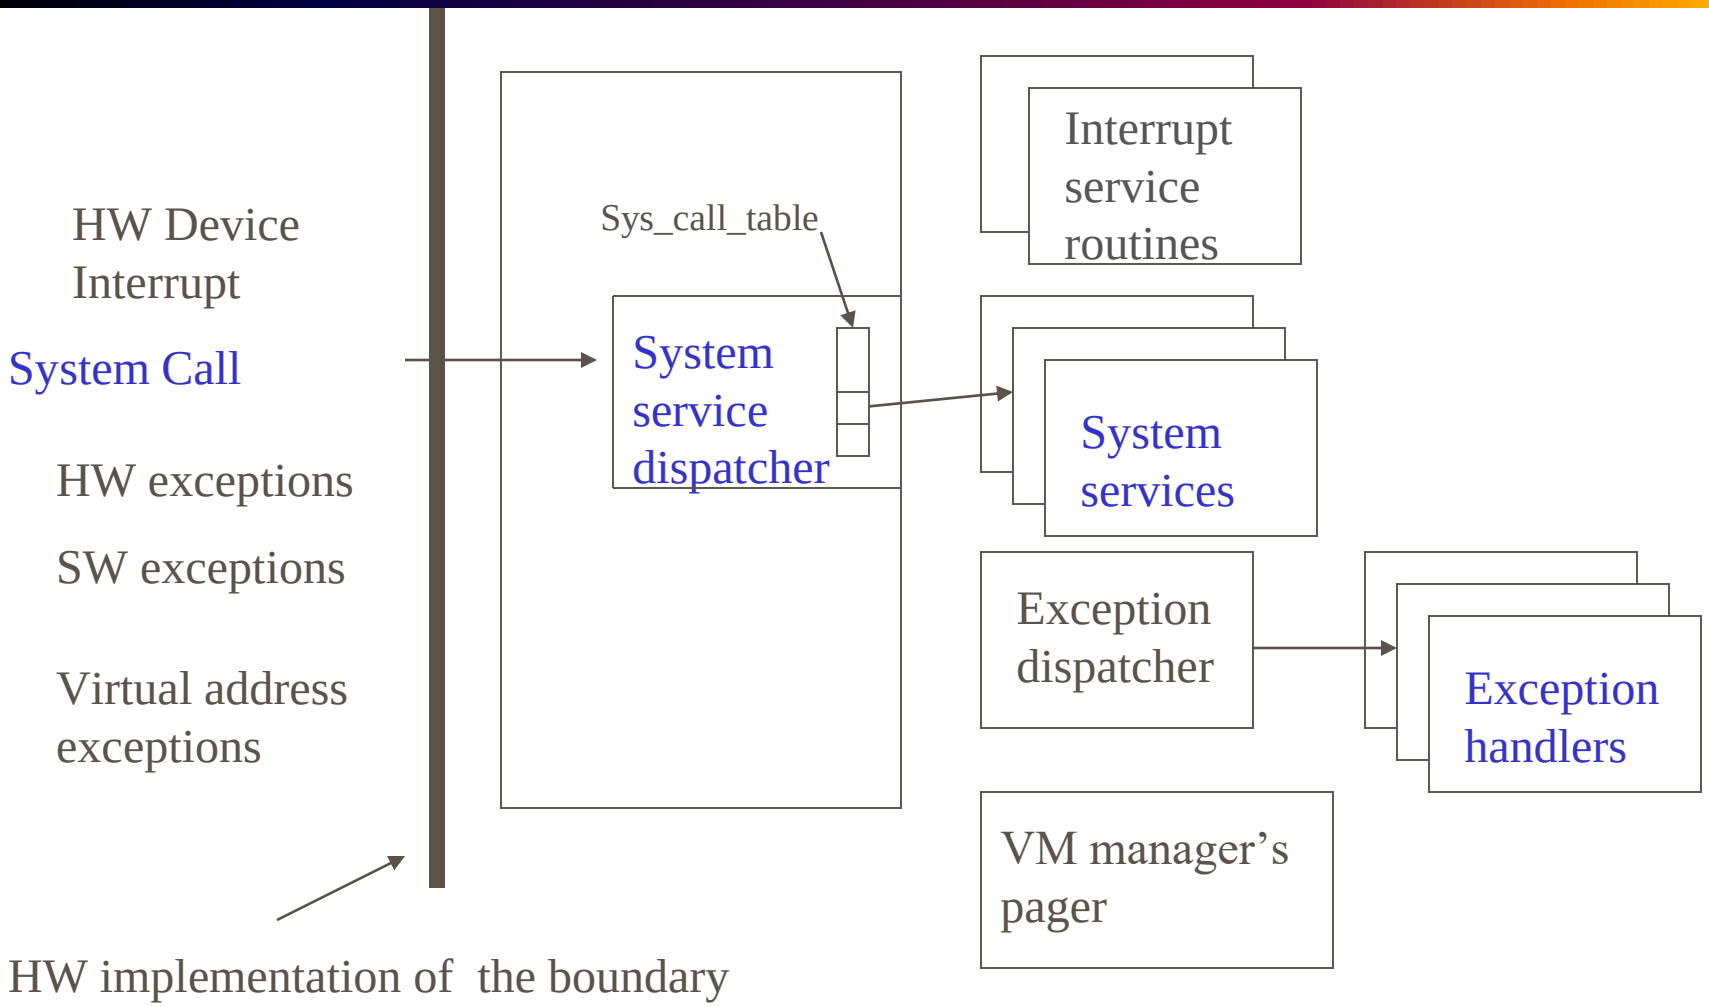

System Call Entry Point

- Assume passing parameters in registers

  EntryPoint:

  switch to kernel stack

  save context

  check register

  call the real code pointed by register

  restore context

  switch to user stack

  iret (change to user mode and return)

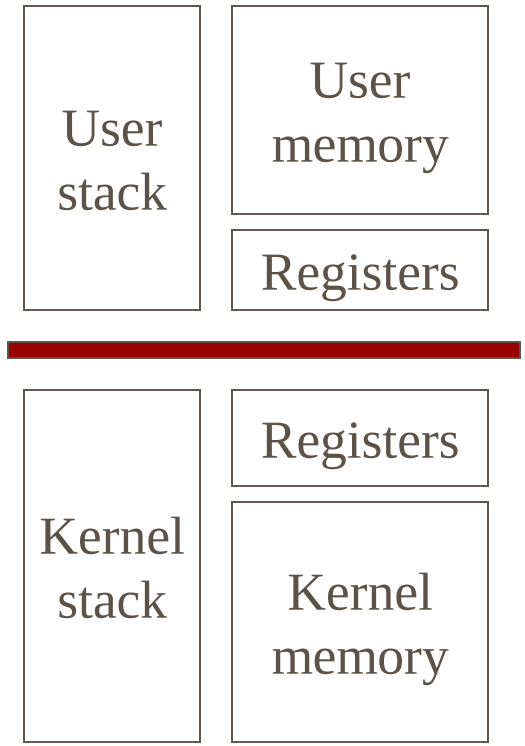

System Call in Linux

- In Linux, all system calls use the following conventions
  - INT 0x80 to invoke (IDT vector 0x80)
  - EAX = system call # (asm/unistd.h)
  - EAX = return value (negative for errors) (asm-generic/errno.h)

System Call

- Vector in IDT is **`system_call`** (in arch/i386/kernel/entry.S)
  - saves registers to stack
  - check for a valid system call #
  - call specific routine using a jump table: **`sys_call_table`** (syscall_table.S, included from entry.S)
  - stack (in kernel level) now appears as shown… (note the argument order)

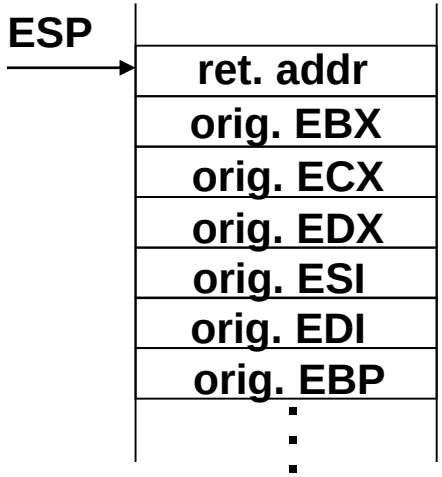

System Call Number

- Each system call is assigned a unique syscall number

```
574 .data
575 ENTRY(sys_call_table)
576         .long sys_restart_syscall       /* 0 - old "setup()" system call, used for restarting */
577         .long sys_exit
578         .long sys_fork
579         .long sys_read
580         .long sys_write
581         .long sys_open          /* 5 */
582         .long sys_close
583         .long sys_waitpid
584         .long sys_creat
585         .long sys_link
586         .long sys_unlink        /* 10 */
587         .long sys_execve
588         .long sys_chdir
589         .long sys_time
590         .long sys_mknod
591         .long sys_chmod         /* 15 */
592         .long sys_lchown16
593         .long sys_ni_syscall    /* old break syscall holder */
594         .long sys_stat
595         .long sys_lseek
596         .long sys_getpid        /* 20 */
597         .long sys_mount
598         .long sys_oldumount
599         .long sys_setuid16
600         .long sys_getuid16
```

The Internal of A Simple System Call

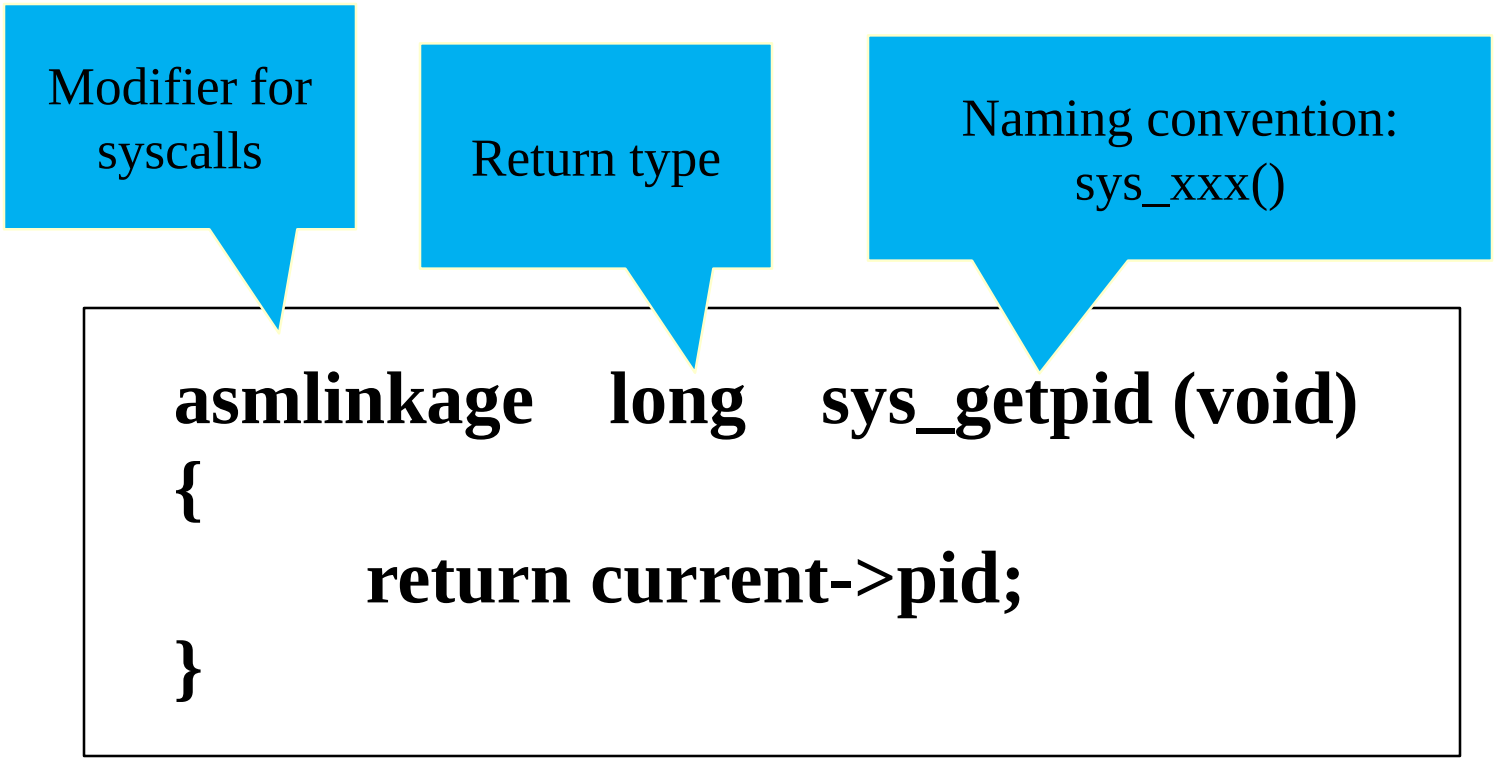

An Example: open()

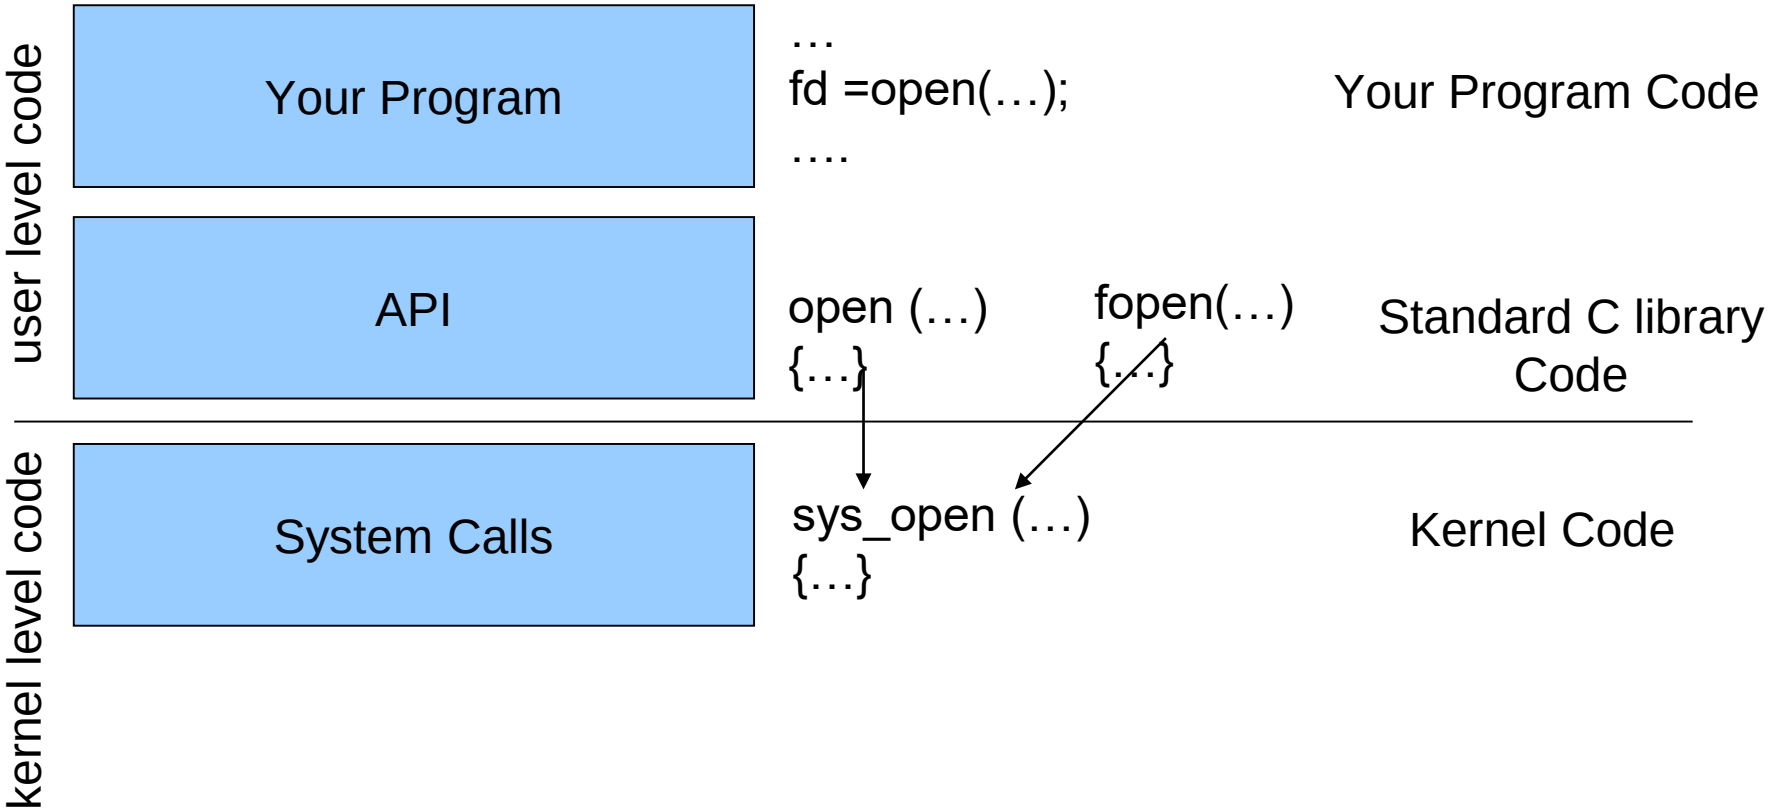

An Example: open()

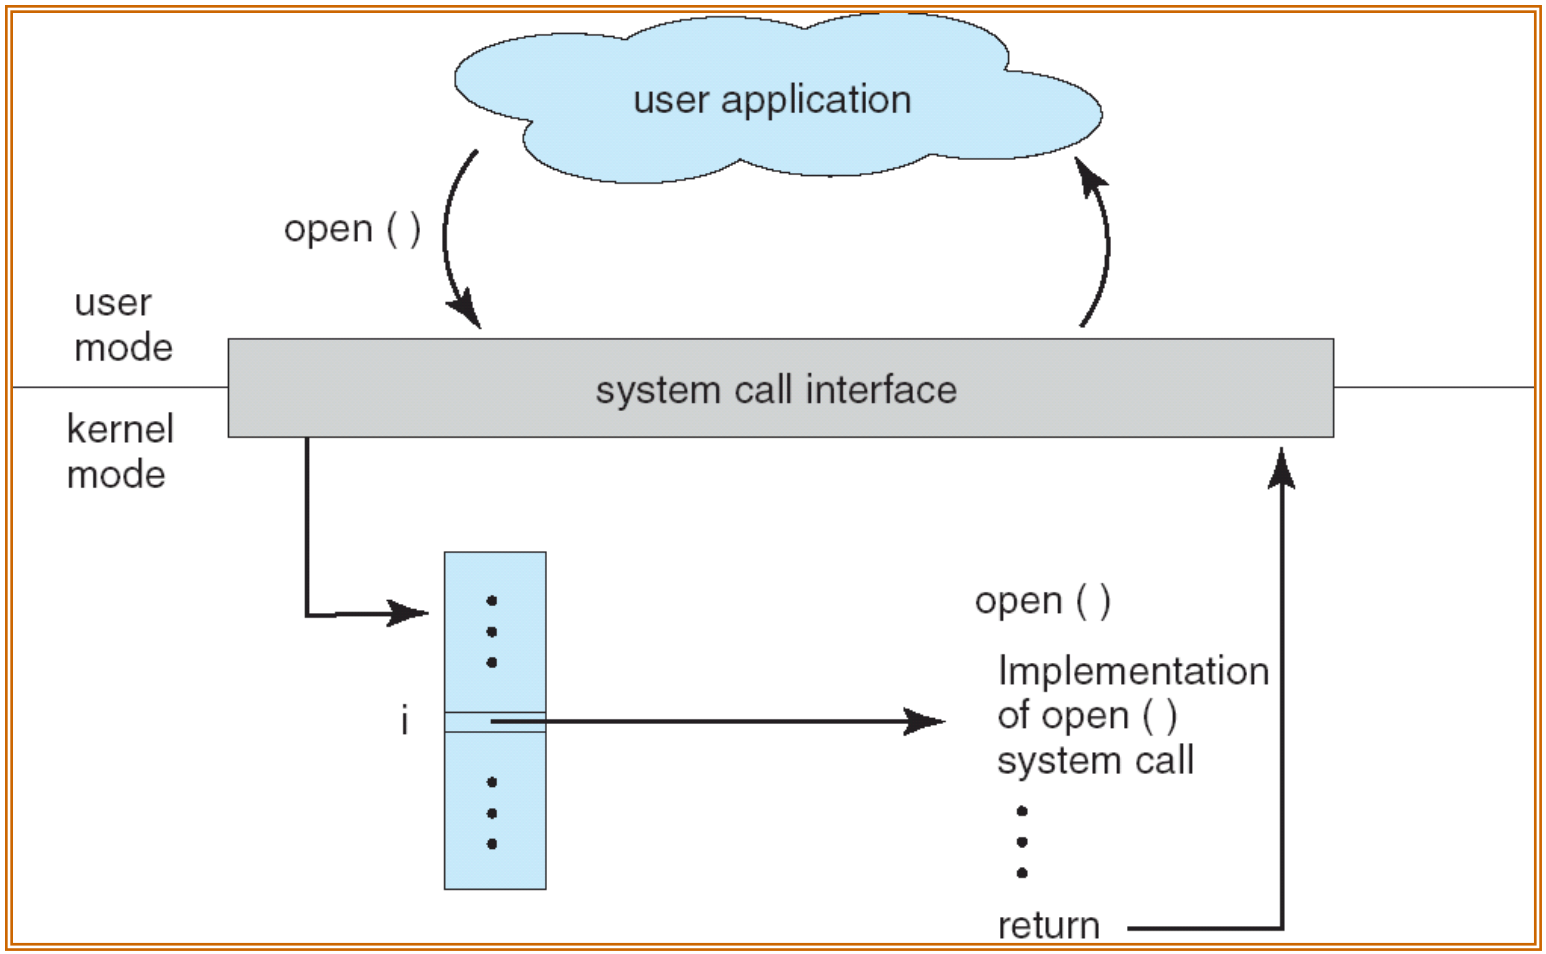

From Lib Code to System Call

- C library code
  - arranges the arguments
  - before performing the system call
  - for example, open and sys_open

```
int open (const char* name, int flags, int mode);

asmlinkage long sys_open (const char* name,
                          int flags, int mode);
```

Open System Call Wrapper

```
open: pushl   %ebx              # save EBX to stack   collee saved register
      movl    0x10(%esp),%edx   # EDX   mode
      movl    0x0C(%esp),%ecx   # ECX   flags
      movl    0x08(%esp),%ebx   # EBX   name
      movl    $0x05,%eax        # open is system call #5
      int     $0x80             # do the system call
      cmpl    $0xFFFFF001,%eax  # -1 to -4095 are errors
      jb      done              # others are valid descriptors
      xorl    %edx,%edx         # negate error number
      subl    %eax,%edx
      pushl   %edx              # save EDX (caller-saved)
      call    __errno_location  # get pointer to errno
      popl    %ecx              # pop error number into ECX
      movl    %ecx,(%eax)       # save error number in errno
      orl     $0xFFFFFFFF,%eax  # return -1
done: popl    %ebx              # restore EBX from stack
      ret
```

Open System Call Wrapper

- Call to **`open`** translates to a call to **`sys_open`** inside the kernel

- Before going on, let's see error handling path
- Recall that, on error, C library
  - returns -1 (for many calls, including **`open`**)
  - stores error code in **`errno`** variable

- But library code is relocatable, so **`errno`** address is not fixed

Open System Call Wrapper

- In dynamically-linked code
  - code addresses generally linked through jump tables
  - data addresses often use the trick
    - make a “fake” call
    - call pushes return address onto stack
    - return address is located at fixed offset from static variable
    - load from stack and add offset to obtain pointer to variable

More on sys_open

- **`sys_open`** (in fs/open.c) invokes **`do_sys_open`**
- **`do_sys_open`** does the following
  - get a free file descriptor
  - open the file (using **`do_filp_open`**)
  - attach the file to the descriptor

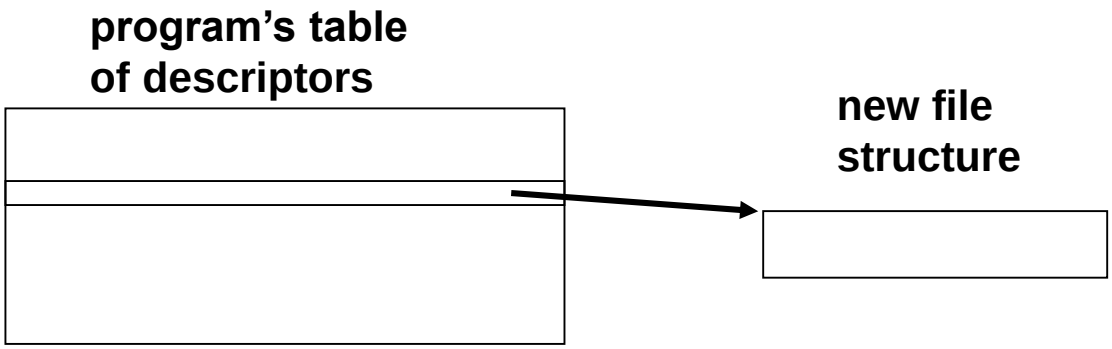

Open System Call Wrapper

- **`filp_open`**
  - check access rights
  - find VFS mount
  - get a file structure from a free list
  - open directory entry (using **`__dentry_open`**, below)

- **`__dentry_open`**
  - fill in file structure
  - if **`fops`** structure is not NULL && **`fops->open`** entry is not NULL, call **`fops->open`**

Interrupt Invocation

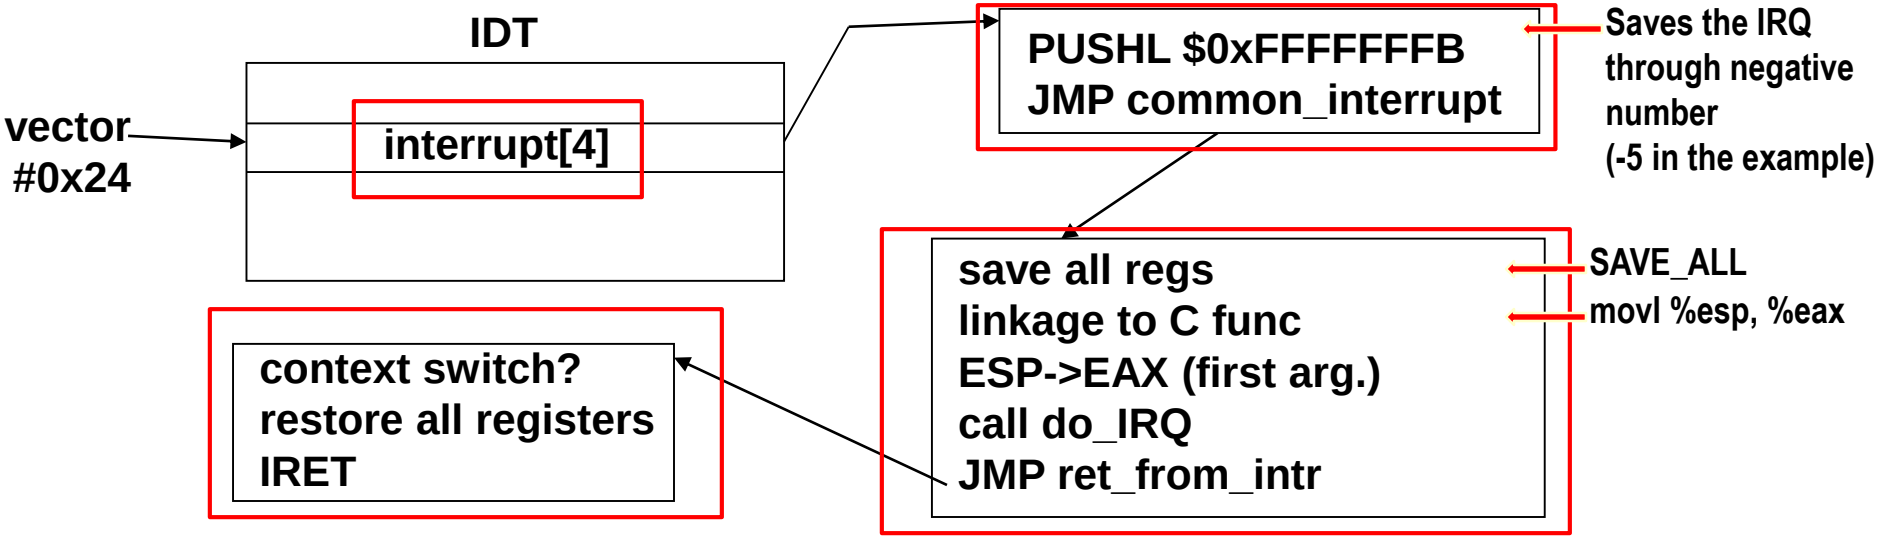

- Why funnel all IRQs into one piece of code(instead of producing one function per IRQ) ?
  - kernel code/size versus speed tradeoff
    - perhaps no big deal for 16
    - keep in mind that you're saving a tiny number of instructions
    - APIC used by most SMPs has 256 vectors

Stack on Entry to do_IRQ