ECE391 Computer System Engineering Lecture 27

Dr. Zbigniew Kalbarczyk
University of Illinois at Urbana- Champaign

Fall 2018

Lecture Topics

• I-mail design (step 5+)

Aministrivia

- MP3.5 Code cutoff:
 - 5:00pm Monday, December 10
- Hand-in/demo
 - 0.5 h (all members should be there)
 - (Monday night); Tuesday; Wednesday
- Post time slots the next week
 - First come first served

- Competition
 - Thursday December 13, in the ECE391 Lab

Aministrivia

- Final Exam:
 - 8:00am to 11:00am, Friday, December 14

- Review session
 - During the Lecture Time om Thursday, December 6

Driver Design Process

- contemplate security (hard/impossible to add in correctly as afterthought!)
- 1. write out all ops (in terms of visible interface)
- 2. design data structures
- 3. pick a locking strategy
- 4. determine locks needed, pick lock ordering
- 5. consider blocking issues & wakeup

Driver Design Process

- 6. consider dynamic allocation issues & hazards
- 7. write code
- 8. write subfunctions and synchronization rules
- 9. return to 3 or 4 if 7 or 8 fails
- 10. unit test

- Consider a task that can't make progress immediately
 - needs a semaphore
 - needs a page currently on disk
- In these cases
 - task puts itself to sleep
 - rather than wasting CPU cycles (as with spin lock)

- What actually happens when a task puts itself to sleep?
 - recall TASK_INTERRUPTIBLE and TASK_UNINTERRUPTIBLE states
 - task puts itself into queue and tells scheduler to run something else
 - some other task eventually wakes it up

Complexity!

- race conditions between
 - task checking sleep condition and putting itself to sleep
 - some other task looking at queue and possibly waking it up
- problems historically
 - led Linux to create macros
 - to keep most programmers from writing the code
- Data structure in Linux: wait queues
 - wait_queue_head_t
 - forms a doubly-linked list of tasks waiting for some event

Waiting for an event

- again, a macro
- condition may be evaluated many times
- returns 0 when condition is true,-ERESTARTSYS if interrupted by signal

```
int ret val = 0;
if (!(condition)) {
    while (1) {
        // start critical section using (wq->lock)
               add to wait queue (protected by lock)
              (memory barrier)
        // set task state to TASK INTERRUPTIBLE
        // end critical section using (wq->lock)
        if (condition) {break;}
        if (!signal pending (current)) {
            schedule (); // sleep
            continue;
        ret val = -ERESTARTSYS;
        break; // deliver signal, then maybe return
                          schedule() call: 1) catches "last-minute" signals and
return ret val;
                          2) removes task from runqueue
```

Waking Tasks up

```
void wake_up (wait_queue_head_t* wq);
```

wake up one task waiting on a wait queue

```
void wake_up_all (wait_queue_head_t* wq);
```

wake up all tasks waiting on a wait queue

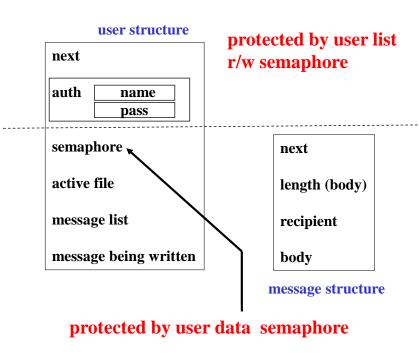
Waking Tasks up

wake up one interruptible task waiting on a wait queue

wake up all interruptible tasks waiting on a wait queue

- What can block?
 - read wait for a new message
 - poll wait for a new message
 - write typically needs to wait for the device, but not in I-mail
 - (so only readers/pollers block in I-mail)
- When are readers (or pollers) awoken?
 - new message delivered
 - user is deleted (nothing is coming now!)

- Data structure impact
 - add a wait_queue_head_t to the user data
 - wait queues
 - have their own synchronization
 - so use it without obtaining user data semaphore
 - (process doing the waking and sleeping processes can't all hold one semaphore simultaneously)



- Protocol for sleeping
 - release all locks
 - check conditions <u>without</u> locks
 - go to sleep
 - when awoken
 - reacquire locks
 - recheck <u>all</u> validity requirements for waking (otherwise go back to sleep; false alarm)

- Q: What implications does check conditions without locks step have?
- A: condition <u>must be</u> atomically safe, i.e., data structure can <u>never</u> be in a state in which evaluation of the condition can lead to a crash

- Important to note
 - sleeper does NOT have locks
 - safe with respect to critical sections is NOT adequate

- Example of things that are ok
 - reading an integer (and comparing it with a constant values)
 - adding two integer fields together
 - reading (but not dereferencing) a pointer and comparing it to NULL

- example of things that might not be OK
 - dereferencing a pointer
 - walking a pointer-based data structure

- in some instances, necessary to restructure
 - critical sections recalculate condition, store as integer
 - code being woken reads integer (always safe)

```
static ssize t
                                                                                       Example I-mail:
Imail read (struct file* f, char* buf, size t len, loff t* offp)
                                /* user data structure */
                                                                                          Imail read()
    Imail user data t* udata;
    Imail msg t* msg;
                                 /* message to be read
    ssize t n read;
                                 /* bytes read
    loff t off;
                                 /* offset
    /* If user has not authenticated yet, permission is denied. */
    if (NULL == (udata = f->private data))
                                                        no lock available in unauthenticated state
        return -EPERM;
                                                                                                           file
                                                                                                           structure
    /* Loop to wait for data to be available for reading. */
                                                                                                       private data
    while (1) {
                                                                                 owned by I-mail
                                                                                                                owned by kernel
        /* Start by obtaining user data lock and checking for user deletion. */
                                                                                   user list
                                                   acquire the lock
        down (&udata->sem):
        if (NULL == udata->active file) {
                                                                                    admin
            up (&udata->sem);
            return -EPERM;
                                                                                                          mailbox
        /* Have data to read? Break out of this loop. */
                                                           if a message is found, leave the loop with the lock
        if (NULL != (msg = udata->msg list))
                                                                                                          writing
            break;
         * No data yet. Release semaphore and either return failure (for
         * non-blocking I/O) or wait for a message or user deletion.
         */
                                          release the lock and go to sleep
        up (&udata->sem);
        if (0 != (f->f flags & O NONBLOCK))
                                                            check asynchronous I/O flag before sleeping
            return - EAGAIN;
        if (wait event interruptible
             (udata->read queue, (NULL == udata->active file ||
                                                                       conditions to wakeup
                                  NULL != udata->msg list)))
            return -ERESTARTSYS; /* got a signal */
         * ERESTARTSYS prevents use of signals to break out of system calls.
         * The semantics are taken from BSD. To change this property, use
         * sigaction and turn off SA RESTART for the signal used to kick
         * system calls out of blocking.
         */
```

Comments on Imail_read()(1)

- Before starting, check the state of file f
 - no lock available in unauthenticated state; throw it out first
 - acquire lock before checking for deletion
- Questions
 - Q: Why not finish state check before acquiring lock?
 - A: prevent race in which user deletion occurs between check and use
 - Q: Why check in the loop?
 - A: locks must be released to sleep, thus user may be deleted between loop iterations (as desired—admin should not have to wait for user program to run!)

Comments on Imail_read()(2)

- Critical section from down to up in loop
 - if a message is found, leave the loop with the lock
 - otherwise, release lock and go to sleep
- Check asynchronous I/O flag before sleeping; if found, return error immediately

- Conditions passed to wait_event_interruptible dereference udata
 - Q: Why is this dereference safe?
 - A: User data is only actually freed in release, which can't be called while some process is in this function.

Comments on Imail_read()(3)

ERESTARTSYS

- What should happen if a signal is delivered while the program sleeps?
- Different operating systems chose different answers (keep going, stop with a specific error, EINTR,).
- Linux allows you to choose.
- ERESTARTSYS options implemented in system call assembly linkage

Example I-mail: Imail read()

```
We have an undeleted user, a message to read, and a semaphore to
   protect us. Note that the default kernel seek functions do not
 * allow file offsets < 0.</pre>
n read = 0:
off = *offp;
                                  check whether there are bytes to read
if (msq->length > off) {
    /* There are bytes left to read--read as many as fit in buffer. */
    if (len < (n read = msq->length - off))
        n read = len;
    /* Copy the bytes into the user's buffer. */
    if (copy to user (buf, msg->body + off, n read))
        n read = -EFAULT;
    else
        (*offp) = off + n read;
 * Release the semaphore and return the number of bytes read
 * (or error number).
up (&udata->sem); ← release the semaphore and return
return n read;
```

Comments on Imail_read()(4)

- At this point, we're ready to read some data
 - user is not deleted
 - user has a message
 - we hold the user data semaphore
- Initialization
 - the default number of bytes read is 0
 - means EOF for read semantics

 Check whether bytes remain in the message (offset vs. length)

Comments on Imail_read()(5)

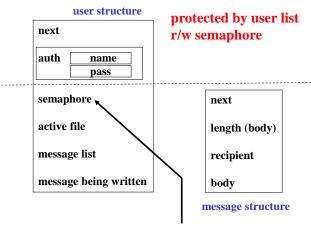
- Make sure that we don't overflow the user's buffer
- Try to read (copy_to_user)
 - copying can cause page faults
 - ok to call while holding semaphores (as in I-mail)
 - not ok to call with spin locks
 - on failure, return EFAULT
 - on success
 - return number of bytes read
 - update the file offset
- In either case, release the semaphore and return

I-mail – Dynamic Allocation

- Questions that we consider
 - does allocation occur inside or outside of critical sections?

udata->writing =

- should we use GFP_KERNEL or GFP_ATOMIC?
- no spin locks in I-mail,
 so GFP_KERNEL is adequate
- Messages
 - allocated by write message operation (ioctl)
 - deleted by
 - delete message
 - I-mail shutdown
 - delete user kfree (udata->msg_list)
 © Steven Lumetta, Zbigniew Kalbarczyk. ECE391
 - failed delivery



protected by user data semaphore

kmalloc (sizeof (Imail msg t), GFP KERNEL);

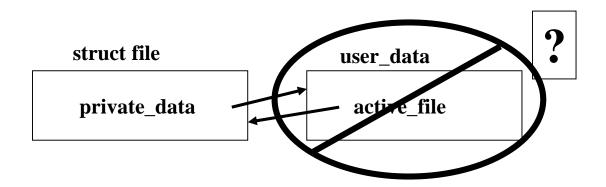
I-mail – Dynamic Allocation

- Users (other than the admin, which is unique static data in the driver)
 - allocated by add user
 - deleted by
 - I-mail shutdown
 - delete user

```
udata = kmalloc (sizeof (Imail_user_data_t), GFP_KERNEL);
kfree (udata)
```

I-mail – Dynamic Allocation

- Dilemma: What if a user is using I-mail when the user data is deleted?
- Can't just delete the user data
 - file structure has pointer
 - may be actively using (changing to NULL may not solve)



Solution to Dilemma

Our solution

- on deletion, remove from user list, but don't free the structure
- give ownership (i.e., responsibility for deletion) to file structure
- it can delete atomically on release call
- signal by changing active_file to NULL

Thus

- if user is not authorized, simply delete
- if user is authorized, set active file to NULL

Few Questions

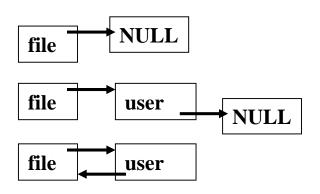
- Q: How do we know that user doesn't authorize after our check?
- A: auth needs read lock, but delete user op holds write lock on user list

- Q: What about after the deletion operation?
- A: user is <u>not in the user list</u> any more!

Few Questions

- Q: Why not be more aggressive? We could try to delete the user data any time that we found the active_file field to be NULL?
- A: Needs to be atomic; basically the same problem as before.
 Release is atomic.

• SO...



unauthenticated (NOT_AUTH)

deleted (DELETED)

authenticated (GOOD)

Sub-operations in I-mail

Find by name (need user list lock of either flavor)

set password

- Used write message (check for recipient in user list)

 - add userdelete userdeliver mail

```
Imail find user (); //(down write (rwsem))
```

Sub-operations in I-mail

Extract message being written (need user data semaphore)

- write message
- Deliver mail (need user list lock, <u>recipient's</u> user data semaphore)
 - release
 - fsync
 - write message

Sub-function Synchronization

- For each sub-function
 - think about interactions with all callers/users
 - select synchronization details on that basis
 - document in function header to avoid misuse

Sub-function Synchronization

- Find by name (Imail_find_user)
 - if user list lock is acquired/released inside of function
 - user may be deleted between search and return
 - thus need to acquire <u>before</u> calling this subfunction
- Extract message being written (Imail_grab_pending)
 - callers must check deletion first
 - deletion might occur if lock is released and reacquired in function
 - thus need to acquire user data sem. <u>before</u> calling this subfunction

Sub-function Synchronization

- Deliver mail (Imail_deliver)
 - allows caller to optionally hold user list lock on entry
 - if not held (as specified by parameter),
 acquires read user list lock
 - acquires recipient's user data semaphore internally
 - releases user list lock if it was acquired within the function

```
static int
Imail delete user (struct file* f, void* arg)
                                                                            Example I-mail:
   Imail user data t* udata;
                             /* user data structure
   int rval;
                             /* return value
                                                                           Imail delete user()
                             /* new user's authentication data */
   Imail auth t auth;
   Imail user data t** find; /* loop index for user removal
   /* Assume success. */
   rval = 0;
   /* This command is only allowed to administrators. */
   if (!Imail user is administrator (f->private data)) ← Check permissions
       return -EPERM;
   /* Read authentication structure into kernel address space.
   if (copy from user (&auth, arg, sizeof (auth)))
       return -EFAULT;
   /* Write lock the user list. */
                                            acquire the lock
   down write (&user list rwsem);
   /* Try to find the user. */
   /* User does not exist. */
       rval = -ENOENT;
                                                             check whether this is administrator
   } else if (udata == &user list) {
       /* Deletion of administrator is not permitted. */
       rval = -EPERM:
   } else {
       /* Discard user's messages and remove them from the user list. */
       Imail user data free (udata);
       for (find = &user list.next; (*find) != udata; find = &(*find)->next);
                                                                            remove the user from the user list
       *find = udata->next;
                                                                           (but don't free it!)
       if (NULL == udata->active file) {
          /* Not active--just delete! */
           kfree (udata);
       } else {
           /* User is active. Boot them off indirectly. */
           udata->active file = NULL;
           wake up interruptible (&udata->read queue);
   /* Release our lock and return the right value. */
   up write (&user list rwsem);
                                             release the lock
   return rval;
```

Comments on Imail_delete_user()(1)

- Start by checking permissions
 - function checks both admin authentication and
 - sysadmin privileges (capability)
- Read the deletion identification before entering the critical section
- Mark the critical section (the rest of the function)

- Check whether such a user exists
- Check whether it's the admin
 - who cannot be deleted

Comments on Imail_delete_user()(2)

- Clean up a bit
 - throw away the mailbox and any message in progress
 - remove the user from the user list (but don't free it!)
 - note that the removal code does not allow for missing users
 - Q: How do we know that another program didn't remove first?
 - A: We have held a write lock on the user list since we located the user.

Comments on Imail_delete_user()(3)

- If the user is not online (authenticated)
 - delete right away
 - no file structure release call to which to delegate deletion anyway
 - again, write lock prevents race with authentication

- If the user is online
 - clear the active file pointer
 - wake up any sleeping user programs